

Lab 2-2. Interprocess Communication – Message Queues

Objectives

- Understand the structure of message queue, and implement the example of server-client program.

Background

1. Name Space:

The table is set of possible names for a given type of IPC. A key is a long int, used to identify individual resources.

| IPC Type | Name space | Identification |
|--------------------------|------------|-------------------|
| PIPE | No name | File descriptor |
| FIFO | Pathname | File descriptor |
| Message Queue | Key | Identifier |
| Semaphore | Key | Identifier |
| Shared Memory | Key | Identifier |
| Socket (Unix Domain) | Pathname | File descriptor |
| Socket (Internet Domain) | Socket | Socket descriptor |

2. Similarities of the 3 types of IPC Channels:

| | Message Queue | Semaphore | Shared Memory |
|--|----------------|-------------|---------------|
| include | <sys/msg.h> | <sys/sem.h> | <sys/shm.h> |
| System calls to create or open | msgget | semget | shmget |
| System call for control operation | msgctl | semctl | shmctl |
| System calls for IPC operations | msgsnd, msgrcv | semop | shmat, shmdt |

3. IPC Flags:

- IPC_CREAT: Creates a unique IPC channel if none exists; otherwise use an existing IPC channel.
- IPC_EXCL: When used with IPC_CREAT, error if the channel exists.


| Flag argument | Key does not exist | Key already exist | Purpose |
|---------------------------------|------------------------|------------------------|------------------------------------|
| NO special flags | Error, errno=ENOENT | OK | to open |
| IPC_CREAT | OK, create new entry | OK, returns same id | To create or to open |
| IPC_CREAT IPC_EXCL | OK, create new entry | Error, errno=EEXIST | To distinguish exist or noexist |

4. Necessary of Message Queue: Drawback of PIPE and FIFO

- 1) Pipe or FIFO is an unformatted stream. Message queue is a formatted stream consisted of messages.
- 2) Pipe or FIFO has to be read in the same order as they are written. Message queue can be accessed randomly.
- 3) Pipe or FIFO is unidirectional. Message queue is bidirectional.
- 4) Pipe or FIFO is simplex. Message queue is multiplex.
- 5) Message queue is faster since it is in kernel.

5. Message Queue

- 1) Also known as mailbox, ports
- 2) The explicit and symmetric naming of processes in direct naming
- 3) Limited modularity since changing the name of a process requires changes elsewhere, i.e., in definitions of other processes



4) P or Q call

send(mbox-id, message)

R calls

receive(mbox-id, message)

- The `msgget()` function returns the message queue identifier associated with key. A message queue identifier and an associated message queue and data structure are created if key is equal to `IPC_PRIVATE`, or key does not already have a message queue identifier associated with it and (`msgflg` & `IPC_CREAT`) is non-zero.

| Function | <code>int msgget (key_t key, int msgflg);</code> |
|--|--|
| Function arguments | |
| key | Specifies the message queue key for which to retrieve the msqid. |
| msgflg | Is a flag that indicates specific message queue conditions and options to implement. |
| Return values | |
| If successful, <code>msgget()</code> returns a message queue identifier. On failure, <code>msgget()</code> returns a value of -1 and sets <code>errno</code> to indicate the error. | |

- The `msgsnd()` function sends a message to the queue associated with message queue identifier `msqid`.

| Function | <code>int msgsnd(int msqid, const void *msgp, size_t msgsz, int msgflg);</code> |
|---|---|
| Function arguments | |
| <i>msqid</i> | Is a unique positive integer, created by <code>msgget()</code> , that identifies a message queue and its associated data structure. |
| <i>msgp</i> | Points to a user-defined buffer. |
| <i>msgsz</i> | Is the length of the message to be sent. |
| <i>msgflg</i> | Specifies the action to be taken if one or more of the following are true |
| Return values | |
| If successful, <code>msgsnd()</code> returns a message queue identifier. On failure, <code>msgsnd()</code> returns a value of -1 and sets <code>errno</code> to indicate the error | |

- The `msgrcv()` function reads a message from the queue associated with the message queue identifier that `msqid` specifies and places it in the user-defined structure that `msgp` points to.

| Function | <code>int msgrcv(int msqid, void *msgp, int msgsz, long msgtyp, int msgflg);</code> |
|---|--|
| Function arguments | |
| <i>msqid</i> | Is a unique positive integer, created by a <code>msgget()</code> call, that identifies a message queue |
| msgflg | Points to a user-defined buffer |
| <i>msgsz</i> | Specifies the size, in bytes, of <code>mtext</code> |
| <i>msgtyp</i> | Specifies the type of message requested (3 types: -, 0, +) |
| Return values | |
| If successful, <code>msgrcv()</code> returns the number of bytes actually placed into <code>mtext</code> . On failure, <code>msgrcv()</code> returns a value of -1, receives no message, and sets <code>errno</code> . | |

6. Semaphore

1) Semaphore

- used for synchronization among processes
- used to synchronize the access of the shared memory segment
- Avoid more than two Processes being Busy waiting at the same time (In User Process, it needs to avoid that more than two processes try to change the same variable)
- Each Semaphore has its own waiting process list.
- Maximum number of Semaphore depends on the number of processes waiting for the corresponding process
- In many cases, you need to synchronize processes
- When they share a resource (shared memory, file descriptor, device, etc.)
- When a process needs to wait for a given event
- Are not used for exchanging large amounts of data as pipes, FIFOs
- Let multiple processes synchronize their operations to synchronize the access to shared memory segments
- Since it provides the resource synchronization between different processes, it must be stored in the kernel

2) Semaphores are positive integers with two methods: UP() and DOWN()

- DOWN():
 - If $\text{sem} \geq 1$ then $\text{sem} = \text{sem} - 1$
 - Otherwise, block the process
- UP():
 - If there are blocked processes, then unblock one of them
 - Otherwise, $\text{sem} = \text{sem} + 1$
- Semaphore operations are atomic

3) **There are generally two usages for semaphores:**

- Mutual exclusion: only one process at a time can be within a section of code
 - Start with $\text{sem} = 1$
 - To enter the mutex section: DOWN(sem)
 - To leave the mutex section: UP(sem)
- Wait for a given event:
 - Start with $\text{sem} = 0$
 - To wait for the event: DOWN(sem)
 - To trigger the event: UP(sem)
 - What happens if the event is triggered before the other process waits for it?

4) Functions and Structures

- semget()

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/sem.h>
```

```
int semget(key_t key, int nsems, int semflag);
```

create semaphore or use the existing semaphore

- nsems : number of semaphore
- semflag : combination of the following constants

- semop()

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/sem.h>

int semop(int semid, struct sembuf **opsptr, unsigned int nops);

struct sembuf {
    ushort sem_num;      /* semaphore # */
    short  sem_op; /* semaphore operation */
    short  sem_flg; /* operation flags */
};
```

execute the arithmetic functions for more than one semaphore

- sem_op :
 - sem_op > 0 : add sem_val on the value of existing semaphore, It mean to release the resources controlled by semaphore
 - sem_op = 0 : the called process wait when the value of semaphore reduces to zero
 - sem_op < 0 : the called process wait when the value of semaphore be greater than or equal to sem_op. It means to allocate the resources.

- semctl()

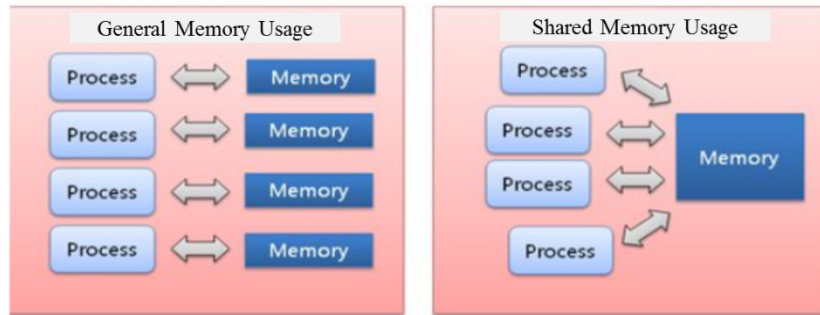
```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/sem.h>

int semctl(int semid, int semnum, int cmd, union semun arg);

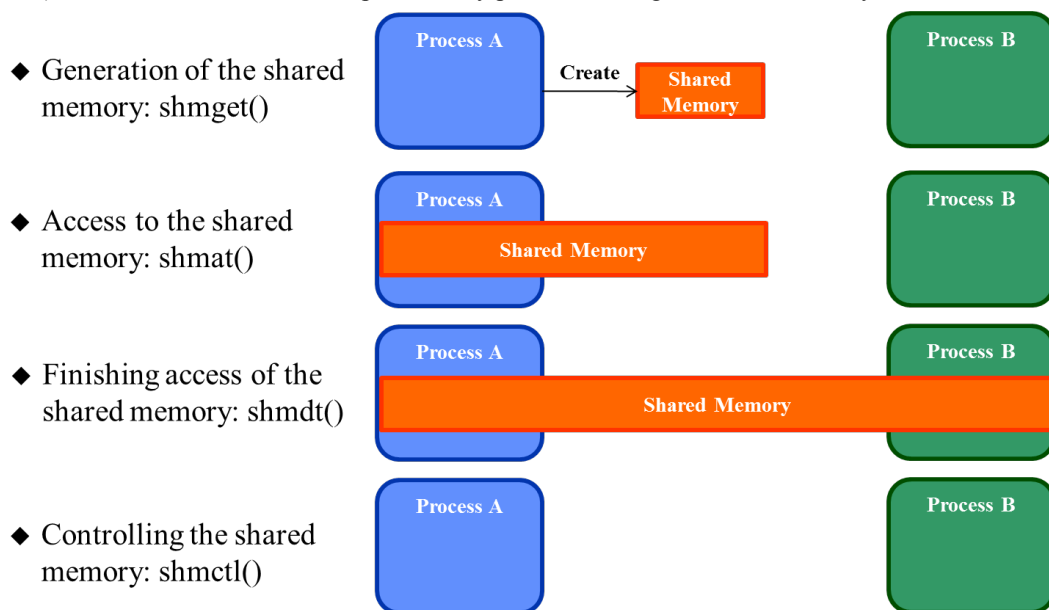
union semun {
    int val; /* used for SETVAL only */
    struct semid_ds *buff; /* used for IPC_STAT and IPC_SET */
    ushort *array; /* used for IPC_GETALL & IPC_SETALL */
} arg;
```

- cmd :
 - IPC_RMID is used to remove semaphore
 - GETVAL is to get the value of semaphore
 - SETVAL is used to change the value of semaphore

7. Shared Memory



- 1) Processes can share the same segment of memory directly when it is mapped into the address space of each sharing process
- 2) Faster communication
- 3) Mutual exclusion must be provided by processes using the shared memory



- The `shmget()` function returns the shared memory identifier associated with key. A shared memory identifier and shared memory segment of at least size bytes is created if key is equal to `IPC_PRIVATE`, or if key does not already have a shared memory identifier associated with it and if (`shmflg & IPC_CREAT`) is non-zero.

| Function | <code>int shmget (key_t key, int size, int shmflg);</code> |
|---|--|
| Function arguments | |
| <i>key</i> | Specifies either <code>IPC_PRIVATE</code> or a unique key. |
| <i>size</i> | Is the size in bytes of the shared memory segment. |
| <i>shmflg</i> | Specifies both the creation and permission bits (for example, <code>IPC_CREAT 0666</code>). |
| Return values | |
| If successful, <code>shmget()</code> returns a shared memory identifier. On failure, it returns -1 and sets <code>errno</code> . | |

- The `shmat()` function attaches the shared memory segment associated with the shared memory identifier *shmid* to the data segment of the calling process.

| Function | <code>void *shmat(int shmid, const void *shmaddr, int shmflg);</code> |
|----------|---|
|----------|---|

| Function arguments | |
|---|--|
| <i>shmid</i> | Is a unique positive integer created by a shmget() system call and associated with a segment of shared memory. |
| <i>shmaddr</i> | Points to the desired address of the shared memory segment. |
| <i>shmflg</i> | Specifies a set of flags that indicate the specific shared memory conditions and options to implement. |
| Return values | |
| If successful, shmat() returns the data segment start address of the attached shared memory segment. On failure, it returns -1, does not attach the shared memory segment, and sets errno | |

- The shmdt() function detaches from the calling process's data segment the shared memory segment located at the address specified by shmaddr.

| Function | int shmdt(const void *shmaddr); |
|--|---|
| Function arguments | |
| <i>shmaddr</i> | Is the data segment start address of a shared memory segment. |
| Return values | |
| If successful, shmdt() decrements the shm_nattach associated with the shared memory segment and returns zero. On failure, it returns -1 and set errno. | |

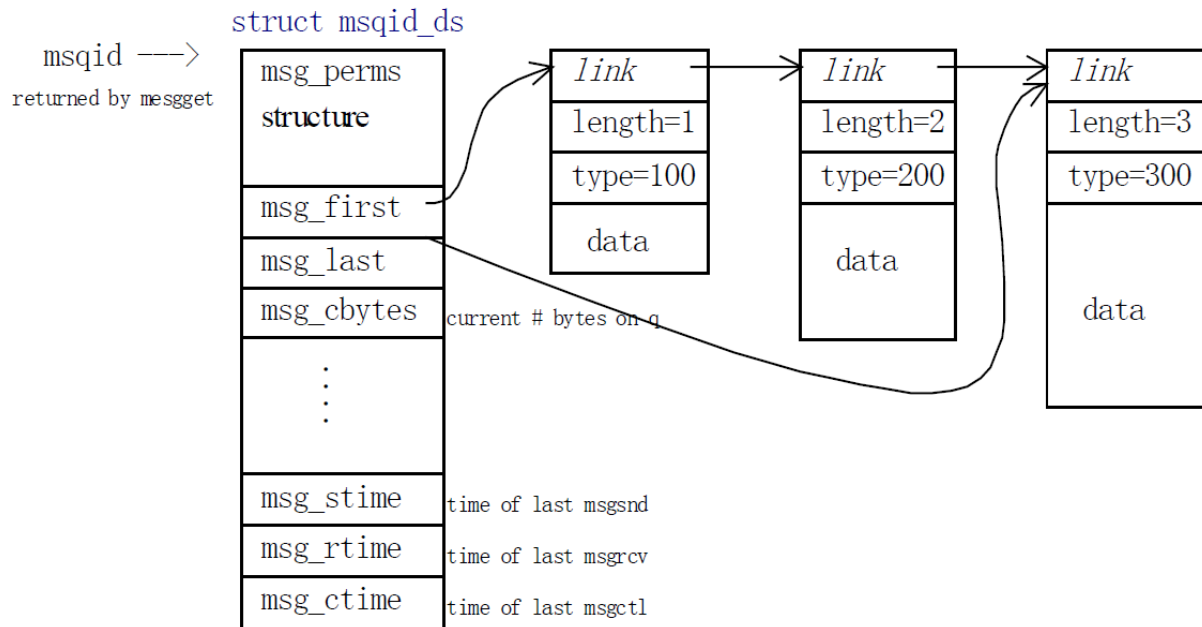
- The shmctl() function provides a variety of shared memory control operations as specified by cmd.

| Function | int shmctl(int shmid, int cmd, struct shmid_ds *buf); |
|---|--|
| Function arguments | |
| <i>shmid</i> | Is a unique positive integer returned by the shmget() function and associated with a segment of memory and a data structure. |
| <i>cmd</i> | Specifies one of IPC_STAT, IPC_SET, or IPC_RMID. |
| <i>buf</i> | Points to the data structure used for sending or receiving data during execution of shared memory control operations. |
| Return values | |
| If successful, shmctl() returns zero. On failure, it returns -1 and sets errno. | |

Practice

Message format: it can be user-defined. Every message has

- size (int), i.e., the length of data,
- type (long int),
- data (if length > 0).



Message Queue Structures in Kernel

```
#define MAXMSGDATA (4096-16)
#define MSGHDRSIZE (sizeof(Msg) - MAXMSGDATA)
typedef struct {
    int msg_len;
    long msg_type;
    char msg_data[MAXMSGDATA];
} Msg;
```

A message queue is created or opened using: **int msgget(key_t key, int msgflag)**

- msgflag: sets permission + IPC_CREAT / IPC_EXCL. See Comments 1 at page 1.
- msgid: returned by the msgget() function, -1 on error.

Messages are sent using: **int msgsnd(int msgid, struct msgbuf *ptr, int length, int flag)**

- struct msgbuf: message format, has a long int (message type) as a field, immediately followed by the data. Must be filled by caller before calling msgsnd().
- length: determine how many bytes of data portion are sent.
- flag: can be IPC_NOWAIT which causes the function to return immediately with error=EAGAIN if no queue space is available. Default action is to block.
- return: 0 if successful; -1 on error.

Messages are received using: **int msgrcv(int msgid, struct msgbuf *ptr, int length, long msgtype, int flag)**

- length: specify the size of the data buffer. Error if the message is too long. OK if the length ≥ message length which is determined by msgsnd(...length...).
- flag: if set to MSG_NOERROR, long messages are truncated and the rest is discarded. If set to IPC_NOWAIT,
- function return -1 if no matching messages are found, otherwise block until

- ▲ A message with the appropriate type is placed on the queue.
- ▲ The message queue is removed from the system.
- ▲ The process receives a signal that is not ignored.
- msgtype = 0: read the first msg on the queue.
 >0: read the first msg on the queue with this type.
 <0: read the first msg on the queue with the lowest type \leq |msgtype|
- return: the length of the data received in bytes if successful; -1 on error.

Control operations use: **int msgctl(int msgid, int cmd, struct msg-id-ts *buff)**

- To remove a msg queue, cmd = IPC_RMID; the 3rd parameter is just a NULL pointer.
- To get/set the queue information in buff, cmd = IPC-STAT/IPC-SET.

Example of message simple client-server using Message Queue:

msg_server.c

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/msg.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include "mesg.h"

#define MKEY1 1234L
#define MKEY2 2345L
#define PERMS 0666

Mesg mesg;

void server(int ipcreadfd, int ipcwritefd) {
    int n, filefd;
    mesg.mesg_type = 1L;
    if((n = mesg_rcv(ipcreadfd, &mesg)) <= 0)
        exit(1);
    mesg.mesg_data[n] = '\0';

    if((filefd = open(mesg.mesg_data, 0)) < 0) {
        sprintf(mesg.mesg_data, "can't open the file\n");
        mesg.mesg_len = strlen(mesg.mesg_data);
        mesg_send(ipcwritefd, &mesg);
    }
    else {
        while((n = read(filefd, mesg.mesg_data, MAXMESGDATA)) > 0) {
            mesg.mesg_len = n;
            mesg_send(ipcwritefd, &mesg);
        }
        close(filefd);
        if(n < 0)
            exit(1);
    }
    mesg.mesg_len = 0;
    mesg_send(ipcwritefd, &mesg);
}

int main(void) {
    int readid, writeid;
```



```

        if((readid = msgget(MKEY1, PERMS | IPC_CREAT)) < 0)
            exit(1);
        if((writeid = msgget(MKEY2, PERMS | IPC_CREAT)) < 0)
            exit(1);

        server(readid, writeid);
        return 0;
    }

```

msg_client.c

```

#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/msg.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include "mesg.h"

#define MKEY1 1234L
#define MKEY2 2345L

Mesg mesg;

void client(int ipcreadfd, int ipcwritefd) {
    int n;
    if(fgets(mesg.mesg_data, MAXMESGDATA, stdin) == NULL)
        exit(1);
    n = strlen(mesg.mesg_data);
    if(mesg.mesg_data[n-1] == '\n')
        --n;
    mesg.mesg_len = n;
    mesg.mesg_type = 1L;
    mesg_send(ipcwritefd, &mesg);

    while((n = mesg_rcv(ipcreadfd, &mesg)) > 0) {
        if(write(1, mesg.mesg_data, n) != n)
            exit(1);
    }
    if(n < 0)
        exit(1);
}

int main(void) {
    int readid, writeid;
    if((writeid = msgget(MKEY1, 0)) < 0)
        printf("client: can't msgget message queue 1\n");
    if((readid = msgget(MKEY2, 0)) < 0)
        printf("client: can't msgget message queue 2\n");

    client(readid, writeid);

    if(msgctl(readid, IPC_RMID, (struct msqid_ds*)0) < 0)
        exit(1);
    if(msgctl(writeid, IPC_RMID, (struct msqid_ds*)0) < 0)
        exit(1);
    return 0;
}

```

mesg.h

```

#include <stdlib.h>
#include <sys/msg.h>

#define MAXMESGDATA (4096-16)
#define MSGHDRSIZE (sizeof(Mesg) - MAXMESGDATA)
typedef struct {
    int mesg_len;
    long mesg_type;
    char mesg_data[MAXMESGDATA];
} Mesg;

void mesg_send(int id, Mesg* mesgptr);
int mesg_rcv(int id, Mesg* mesgptr);

mesg.c

#include "mesg.h"

void mesg_send(int id, Mesg* mesgptr) {
    if(msgsnd(id, (char*)&(mesgptr->mesg_type), mesgptr->mesg_len, 0) != 0)
        exit(1);
}

int mesg_rcv(int id, Mesg* mesgptr) {
    int n;
    n = msgrcv(id, (char*)&(mesgptr->mesg_type), MAXMESGDATA, mesgptr->mesg_type,
0);
    if((mesgptr->mesg_len = n) < 0)
        exit(1);
    return n;
}

```

Result:

```

ancl@cloud03:~/lab$ cat msg_queue_test.txt
message queue test
interprocess communication test

practice
example

ancl@cloud03:~/lab$ ./msg_server
ancl@cloud03:~/lab$

ancl@cloud03:~/lab$ ./msg_client
/home/ancl/lab/msg_queue_test.txt
message queue test
interprocess communication test

practice
example

ancl@cloud03:~/lab$

```

Multiplexing Message and bi-directional message queue:

- Only one queue, but used for 2 directions, and multiple clients can read/write the queue.

