

# Lab 12. Building an OpenStack Cloud Platform

**Chan-Hyun YOUN**

**Dept of Electrical Engineering, KAIST**

# LAB 12-1) KVM HYPERVISOR

# KVM Hypervisor



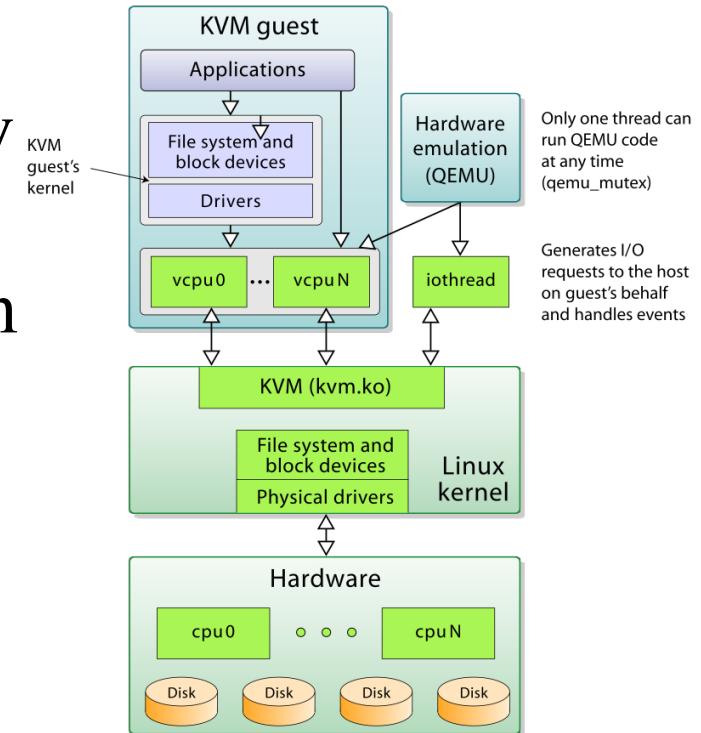
- Kernel-based Virtual Machine (KVM)
  - KVM is a virtualization infrastructure for the Linux kernel that turns it into a hypervisor
  - It was merged into the Linux kernel mainline in kernel version 2.6.20
  - It requires a processor with hardware virtualization extensions

# KVM Hypervisor



- Internals

- Set up the guest VM's address space. The host must also supply a firmware image that the guest can use to bootstrap into its main OS.
- Feed the guest simulated I/O.
- Map the guest's video display back onto the host.
- QEMU support host emulation

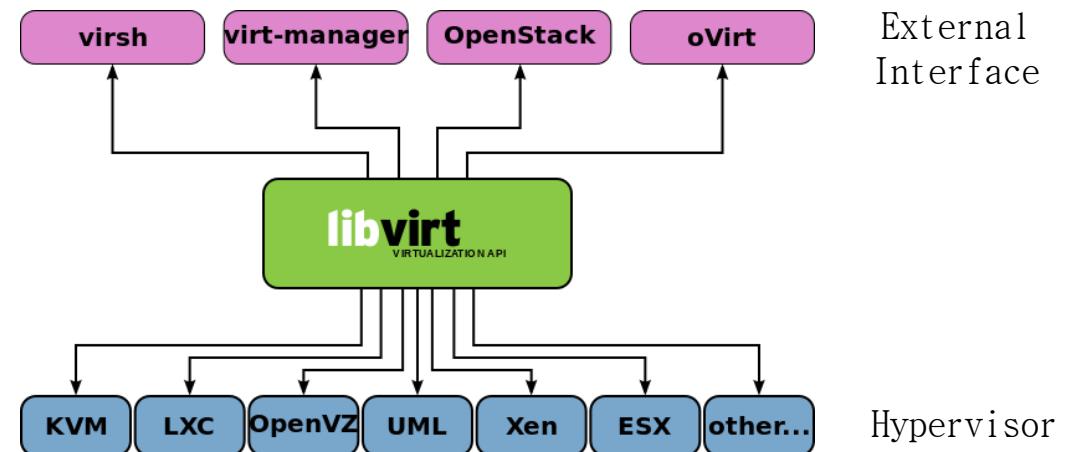


# Libvirt – Hypervisor Orchestrator

- Libvirt



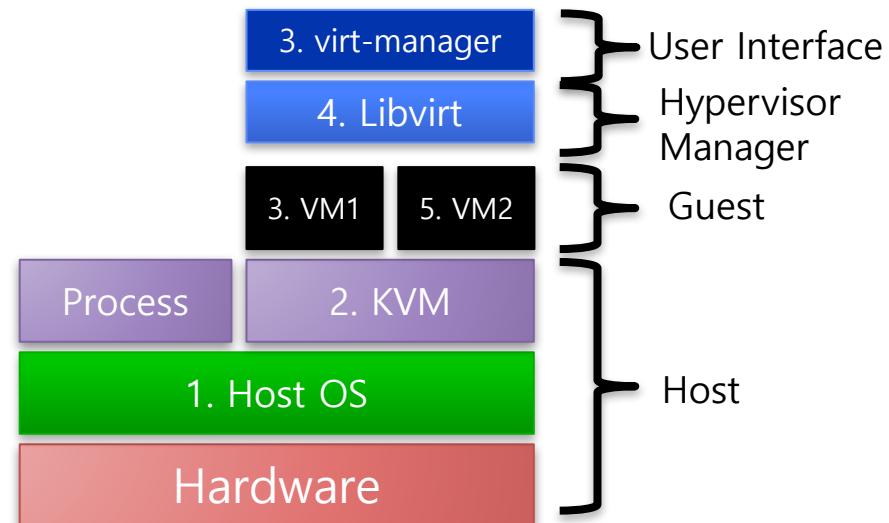
- Open-source API, daemon and management tool for managing platform virtualization
- It can be used to manage KVM, Xen, VMware ESX, QEMU and other virtualization technologies



<https://libvirt.org/>

# Overview of Lab 12-1

1. Prepare Host OS on your computing environment  
(Bare metal PC or VirtualBox)
2. Operate KVM Hypervisor to your Host OS
3. Create Virtual Machine with  
user interface ‘virt-manager’
4. Test the performance  
and manage the VM  
resources via virt-manager  
Virt-manager is operated  
with calling (libvirt api)



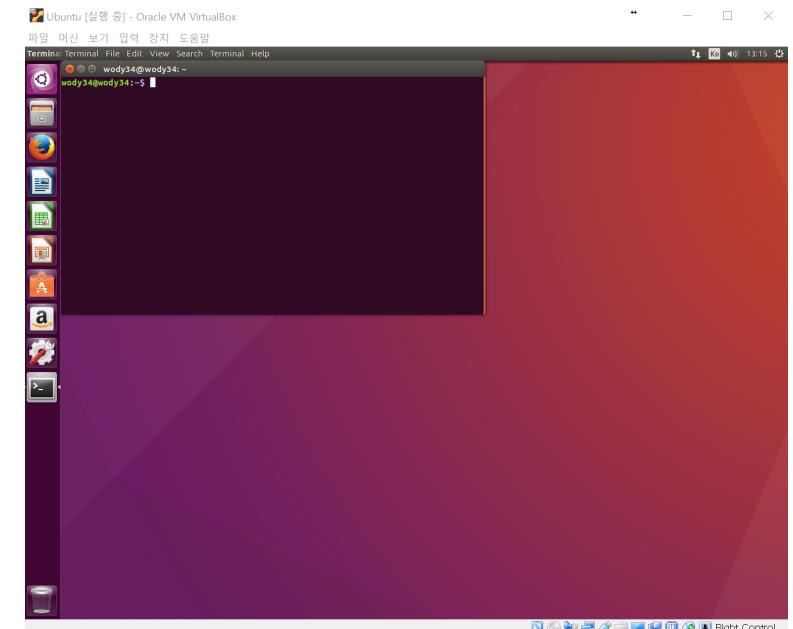
# 1. Prepare Host OS

- Prepare Host OS (Linux Ubuntu 16.04.3 LTS)
  - Your CPU should support **Hardware-assisted virtualization**
    - Intel VT-x or AMD-V
    - If not, the VM is launched under emulated mode (very slow)
    - You can check your CPU information at /proc/cpuinfo

```
processor      : 0                               • Intel: vmx(Hardware virtualization)
vendor_id     : AuthenticAMD                  • AMD: svm(Secure virtual machine)
cpu family    : 16
model         : 4
model name    : AMD Phenom(tm) II X4 925 Processor
cpu MHz       : 2731.284
cache size    : 512 KB
...
flags          : fpu vme de pse tsc msr pae mce cx8 apic sep mtrr pge mca cmov pat pse36
clflush mmx fxsr sse sse2 ht syscall nx fxsr_opt rdtscp lm 3dnowext 3dnow rep_good nopl
extd_apicid  dni lahf_lm cmp_legacy cr8_legacy svm
...
```

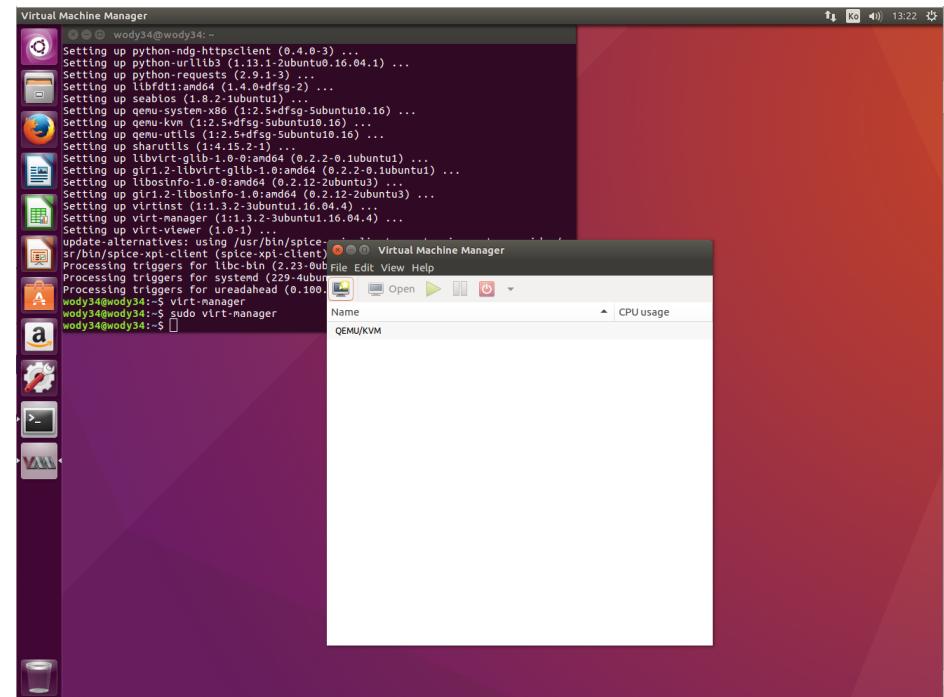
# 1. Prepare Host OS

- Prepare Ubuntu 16.04.3 LTS Desktop version in your computing environment
    - There is way to install Ubuntu in VirtualBox Environment
    - Refer [Ubuntu Installation \(Virtual Box\).pdf](#) in Lab2 supplement
  - Update packages
    - # sudo apt-get update -y
    - # sudo apt-get upgrade -y
    - # sudo apt-get dist-upgrade -y
    - # sudo apt-get install openssh-server git -y



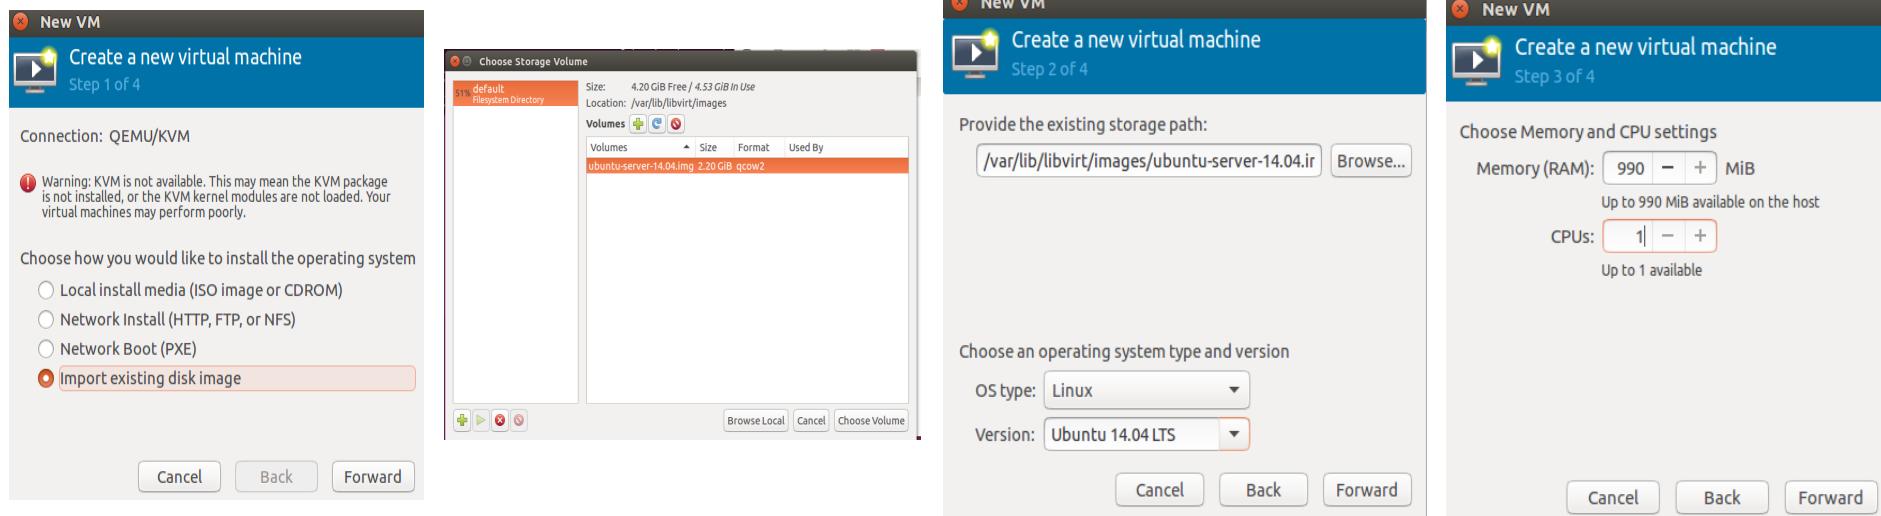
# 2. Install KVM and libraries

- Prepare KVM Hypervisor and virt-manager
  - # sudo apt-get install qemu-kvm libvirt-bin virt-manager
- Run virt-manager
  - # sudo virt-manager



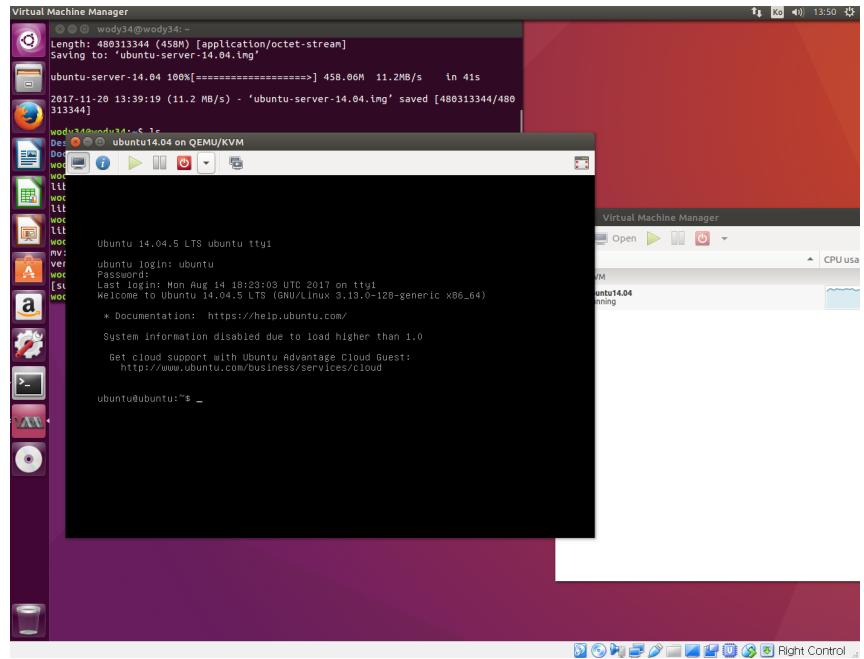
# 3. Create Virtual Machine

- Prepare VM image
  - Download Ubuntu image to `/var/lib/libvirt/images`  
`$ sudo wget http://143.248.147.118/~djshin/ubuntu-server-14.04.img`
  - Click button ‘Create a new virtual machine’



# 3. Create Virtual Machine

- Log-in to Guest OS
  - ubuntu / ubuntu
- Set up root password
  - # sudo passwd root
  - Type Your Password!
- Log-in with Root account
  - # su – root
  - Type Your Password!
- Check Guest's IP address
  - # ifconfig
- Power-off your guest

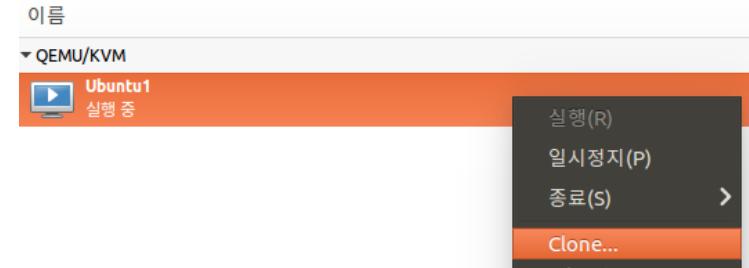


```
ubuntu@ubuntu:~$ su - root
Password:
root@ubuntu:~# _
```

```
root@ubuntu:~# ifconfig
eth0      Link encap:Ethernet  HWaddr
          inet  addr:192.168.122.185
```

# 3. Create Virtual Machine

- Clone for 2<sup>nd</sup> VM
  - Clone : Make a new copied VM
  - Right click -> Clone
- Boot both of VMs
  - Check its IP address
    - # ifconfig



The image displays two terminal windows side-by-side. The left window is titled 'Ubuntu1 on QEMU/KVM' and the right window is titled 'Ubuntu2 on QEMU/KVM'. Both terminals show the root prompt and the output of the 'ifconfig' command.

```
Ubuntu1 on QEMU/KVM
http://www.ubuntu.com/business/services/cloud

root@ubuntu:~$ su - root
Password:
root@ubuntu:~# ifconfig
eth0      Link encap:Ethernet HWaddr 52:54:00:44:7c:c3
          inet addr:192.168.122.185  Bcast:192.168.122.255  Mask:255.255.255.0
          inet6 addr: fe80::5054:ff:fe44:7cc8/64 Scope:Link
             UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
             RX packets:27 errors:0 dropped:0 overruns:0 frame:0
             TX packets:40 errors:0 dropped:0 overruns:0 carrier:0
             collisions:0 txqueuelen:1000
             RX bytes:3474 (3.4 KB)  TX bytes:3872 (3.8 KB)

lo       Link encap:Local Loopback
          inet addr:127.0.0.1  Mask:255.0.0.0
          inet6 addr: ::1/128 Scope:Host
             UP LOOPBACK RUNNING  MTU:65536  Metric:1
             RX packets:0 errors:0 dropped:0 overruns:0 frame:0
             TX packets:0 errors:0 dropped:0 overruns:0 carrier:0
             collisions:0 txqueuelen:0
             RX bytes:0 (0.0 B)  TX bytes:0 (0.0 B)

root@ubuntu:~# 

Ubuntu2 on QEMU/KVM

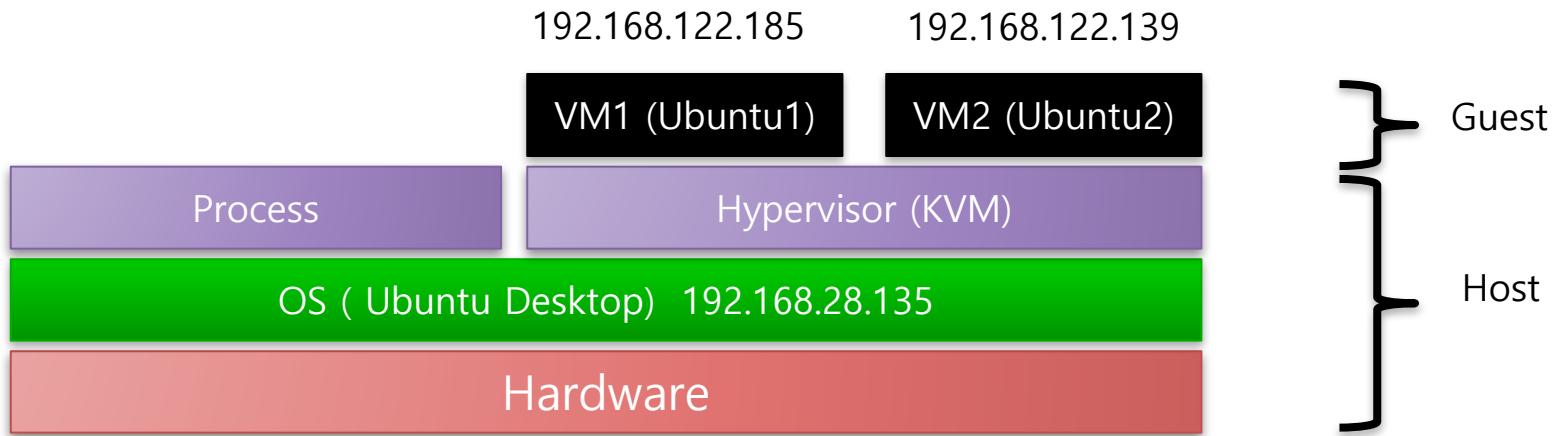
root@ubuntu:~# 
root@ubuntu:~# ls
root@ubuntu:~# 
root@ubuntu:~# ifconfig
eth0      Link encap:Ethernet HWaddr 52:54:00:af:a8:e2
          inet addr:192.168.122.139  Bcast:192.168.122.255  Mask:255.255.255.0
          inet6 addr: fe80::5054:ff:fea8e2/64 Scope:Link
             UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
             RX packets:29 errors:0 dropped:0 overruns:0 frame:0
             TX packets:37 errors:0 dropped:0 overruns:0 carrier:0
             collisions:0 txqueuelen:1000
             RX bytes:2368 (2.3 KB)  TX bytes:3602 (3.6 KB)

lo       Link encap:Local Loopback
          inet addr:127.0.0.1  Mask:255.0.0.0
          inet6 addr: ::1/128 Scope:Host
             UP LOOPBACK RUNNING  MTU:65536  Metric:1
             RX packets:0 errors:0 dropped:0 overruns:0 frame:0
             TX packets:0 errors:0 dropped:0 overruns:0 carrier:0
             collisions:0 txqueuelen:0
             RX bytes:0 (0.0 B)  TX bytes:0 (0.0 B)

root@ubuntu:~# ifconfig
```

# 3. Create Virtual Machine

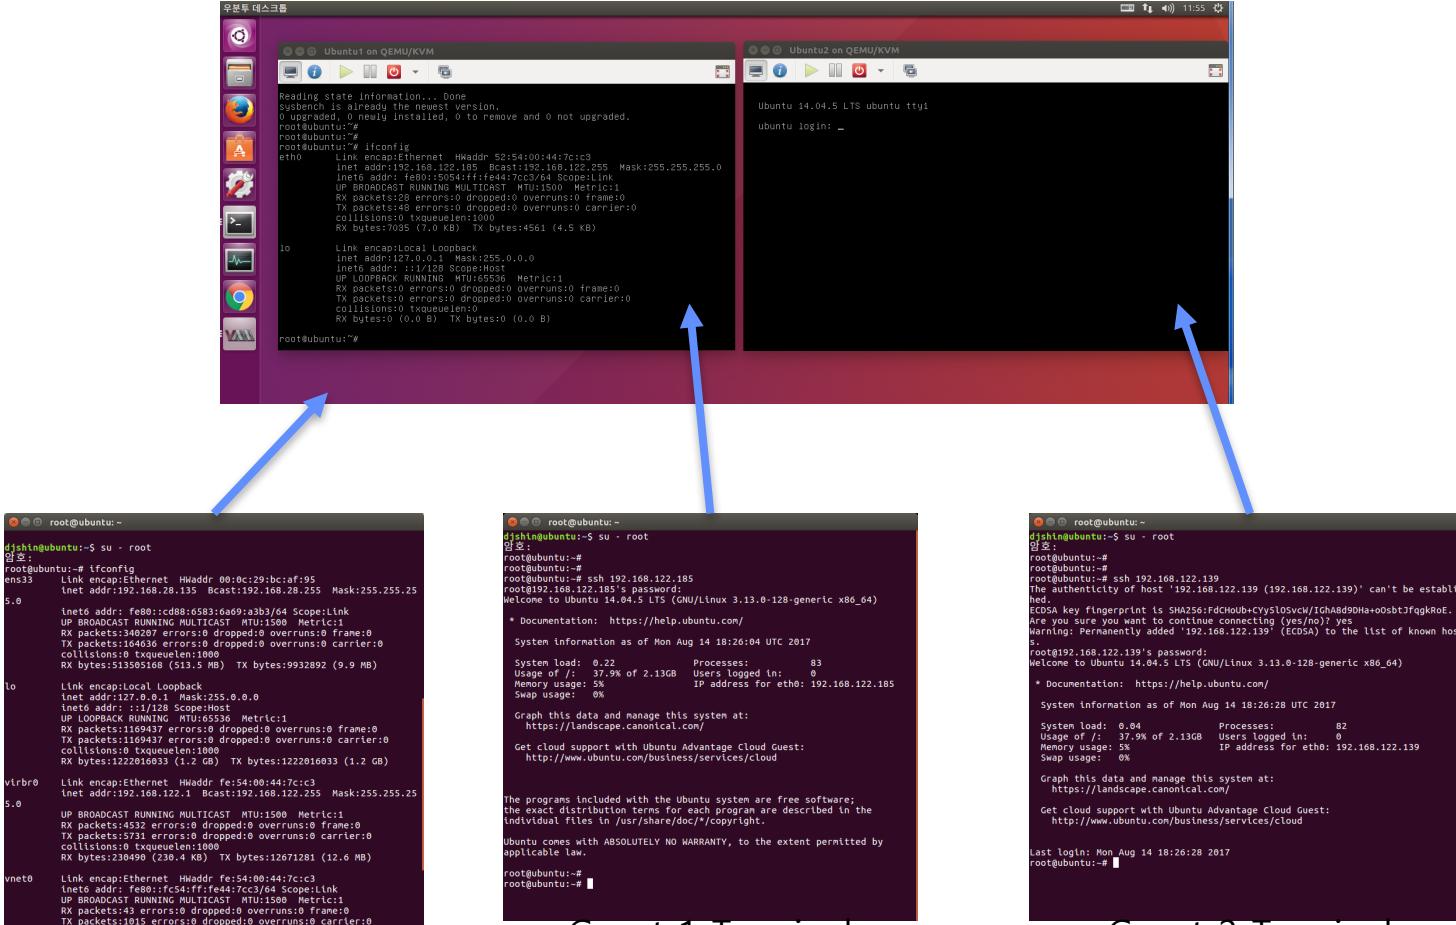
- Configured Cloud Environment



- IP address
  - # ifconfig
- CPU Info
  - # cat /proc/cpuinfo
- Process Info
  - # top
- Free memory
  - # free -m
- Disk Info
  - # df -hT

# 3. Create Virtual Machine

- 3 Terminals for Host, Guest1 , Guest2



Host Terminal

Korea Advanced Institute of  
Science and Technology

Guest 1 Terminal  
(192.168.122.185)

Guest 2 Terminal  
(192.168.122.139)

# 3. Create Virtual Machine

- Host's View

OS ( Ubuntu Desktop) 192.168.28.135

## Hardware

- 6 CPU Cores
  - 8GB of RAM
  - 52GB of Disk  
  - 2 VMs as a Process

```
root@ubuntu:~# cat /proc/cpuinfo | grep processor
processor       : 0
processor       : 1
processor       : 2
processor       : 3
processor       : 4
processor       : 5
```

```
root@ubuntu:~# free -m
```

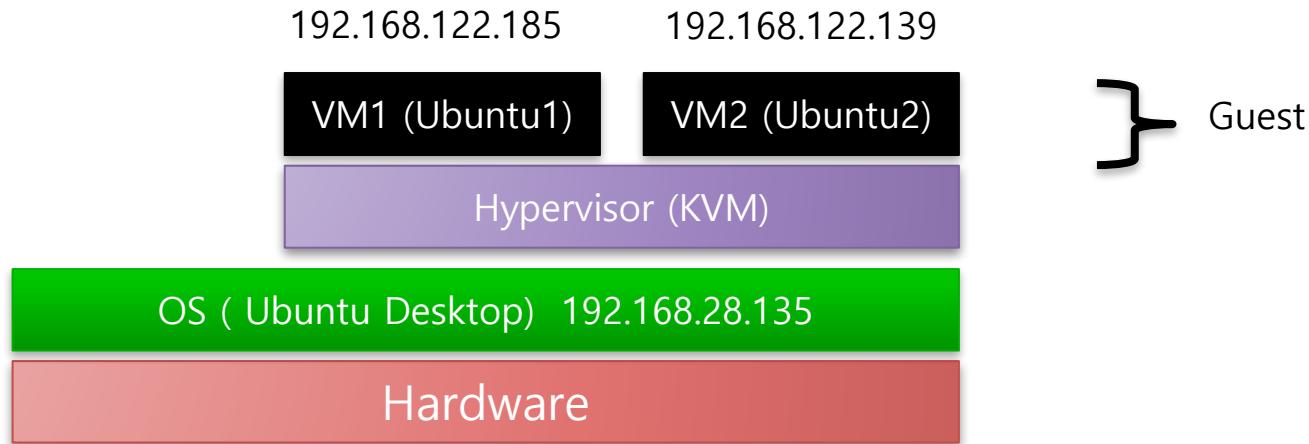
	total	used	free
Mem:	7963	2037	1328
Swap:	8189	0	8189

crypt#	crypt#	Filesystem	Size	Used	Avail	Type
/dev/sda1		ext4	52G	17G	32G	35% /

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
11447	libvirt+	20	0	1870744	307608	18836	S	122.6	3.8	0:52.92	qemu-system-x86
11301	libvirt+	20	0	1965864	318832	19116	S	113.6	3.9	0:53.79	qemu-system-x86

# 3. Create Virtual Machine

- Guest's View



- 2 CPU Cores
- 1GB of RAM
- 2.2GB of Disk

```
root@ubuntu:/# cat /proc/cpuinfo |grep processor
processor       : 0
processor       : 1
```

	total	used	free
Mem:	993	160	832
-/+ buffers/cache:		65	928
Swap:	0	0	0

/dev/sda1	ext4	2.2G	842M	1.2G	41%	/
-----------	------	------	------	------	-----	---

# 3. Create Virtual Machine

- Benchmark
  - A number of standard test (computer program) to assess the relative performance
- We'll use 2 benchmarks
  - Sysbench, stress-ng
  - For each host and guests
  - # apt-get install stress-ng sysbench

# 3. Create Virtual Machine

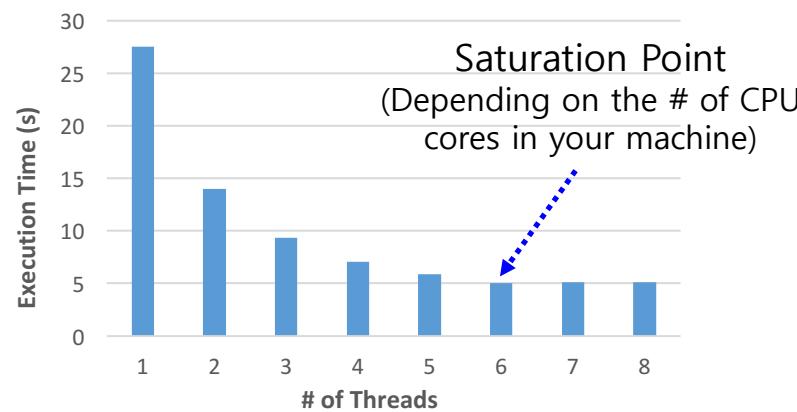
- sysbench CPU Test in Host

- Monitor the result of the test and CPU usage with “top” command
- # sysbench --test=cpu --num-threads=1 --cpu-max-prime=25000 run

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
4934	djshin	20	0	28480	2364	2016	S	100.0	0.0	0:23.87	sysbench

- increasing # of threads
- # sysbench --test=cpu --num-threads=2 --cpu-max-prime=25000 run
- ...
- # sysbench --test=cpu --num-threads=8 --cpu-max-prime=25000 run

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
5055	djshin	20	0	28512	2384	2040	S	200.0	0.0	0:17.24	sysbench



# 4. Virtual Machine Management

- Lifecycle Management
  - VM creation
  - VM deletion
  - VM migration
  - VM snapshot management (save & restore)
- Resource Management
  - CPU Management
    - Change number of assigned cores
  - Memory Management
    - Give or Take Memory
  - Virtual Disk Management
    - Create new disk image and mount

# 4. Virtual Machine Management

- In your host terminal
  - Host node info : # virsh nodeinfo
  - VM list : # virsh list
  - VM info : # virsh dominfo <VM name>

```
root@ubuntu:~# virsh nodeinfo
CPU 모델:          x86_64
CPU:                 6
CPU 주파수:        3500 MHz
CPU 소켓:          2
소켓당 코어:       3
코어당 스레드:     1
NUMA cell:          1
메모리 용량:      8154208 KiB
```

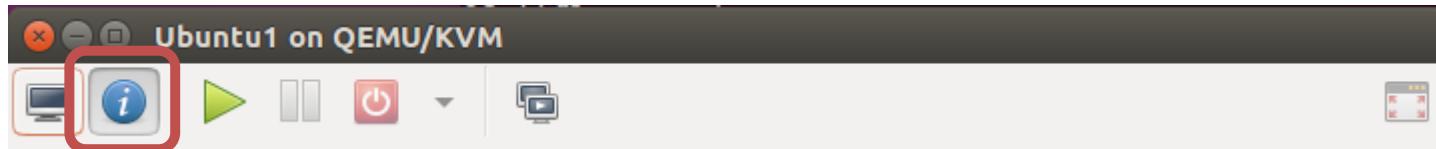
Id	이름	상태
1	Ubuntu1	실행 중
2	Ubuntu2	실행 중

```
root@ubuntu:~# virsh dominfo Ubuntu1
Id:          1
이름:        Ubuntu1
UUID:        0906832b-e51c-4a5c-b76d-d2dbaa8dba7
OS 유형:    hvm
상태:        실행 중
CPU:          2
CPU 시간:   74.8s
최대 메모리: 1048576 KiB
사용된 메모리: 1048576 KiB
Persistent:  예
Autostart:   사용 안 함
Managed save: 아니요
모안 모델:  apparmor
보안 DOI:    0
보안 레이블: libvirt-0906832b-e51c-4a5c-b76d-d2dbaa8dba7 (enforcing)
```

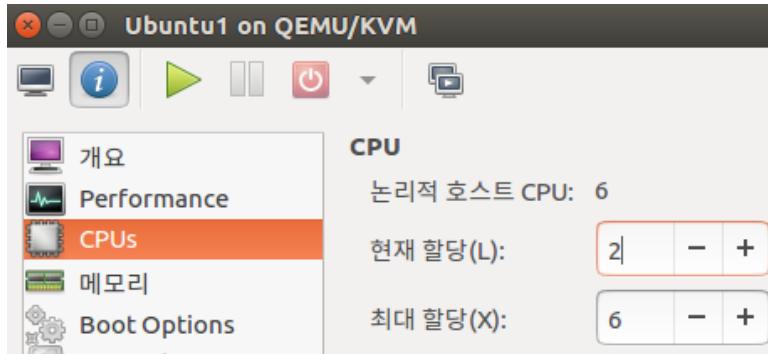
```
root@ubuntu:~# virsh dominfo Ubuntu2
Id:          2
이름:        Ubuntu2
UUID:        bf1fbcd-57d4-42c1-a667-73c9a615d954
OS 유형:    hvm
상태:        실행 중
CPU:          2
CPU 시간:   31.0s
최대 메모리: 1048576 KiB
사용된 메모리: 1048576 KiB
Persistent:  예
Autostart:   사용 안 함
Managed save: 아니요
모안 모델:  apparmor
보안 DOI:    0
보안 레이블: libvirt-bf1fbcd-57d4-42c1-a667-73c9a615d954 (enforcing)
```

# 4. Virtual Machine Management

- Turn off your VMs
- virt-manager -> VM open -> VM Info



- Set each of VMs
  - CPU 2, Memory 1G



- Turn on all VMs

# 4. Virtual Machine Management

- In your guest terminal

- Run CPU Benchmark

- # sysbench --test=cpu --num-threads=1 --cpu-max-prime=25000 run
    - ...
    - # sysbench --test=cpu --num-threads=6 --cpu-max-prime=25000 run

```
Test execution summary:
total time: 14.1875s
```

- In your host terminal

- Give more cores to VM1 (2 to 6)
  - # virsh setvcpus Ubuntu1 --count 6

- In your guest terminal

- Turn on the cores
  - # echo 1 > /sys/devices/system/cpu/cpu2/online
  - # echo 1 > /sys/devices/system/cpu/cpu3/online
  - # echo 1 > /sys/devices/system/cpu/cpu4/online
  - # echo 1 > /sys/devices/system/cpu/cpu5/online
  - # lscpu

```
root@ubuntu:~# lscpu
Architecture:          x86_64
CPU op-mode(s):        32-bit, 64-bit
Byte Order:            Little Endian
CPU(s):                6
On-line CPU(s) list:  0-5
```

- Run benchmark again in the guest VM

- # sysbench --test=cpu --num-threads=6 --cpu-max-prime=25000 run

```
Test execution summary:
total time: 5.2964s
```

# 4. Virtual Machine Management

- In your guest terminal

- Check the size of free memory
  - # free -m
- Run Memory Benchmark
  - # stress-ng --vm 1 --vm-bytes 1G --timeout 20s --metrics-brief
- Cannot run the benchmark due to the out of memory .

```
root@ubuntu:~# free -m
              total        used        free      shared  buffers  cached
Mem:       880         123        757          0         11        50
-/+ buffers/cache:       61        819
Swap:          0          0          0
```

- In your host terminal

- Give more memory to VM1
  - # virsh setmem Ubuntu1 --size 4G

```
root@ubuntu:~# stress-ng --vm 1 --vm-bytes 1G --timeout 20s --metrics-brief
stress-ng: info: [1237] dispatching hogs: 1 vm
stress-ng: info: [1237] successful run completed in 20.00s
stress-ng: info: [1237] stressor      bogo ops real time  usr time  sys time
  ogo ops/s  bogo ops/s
stress-ng: info: [1237]                                     (secs)  (secs)  (secs)
  eal time) (usr+sys time)
```

- In your guest terminal

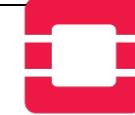
- Check the size of free memory
  - # free -m
- Run Memory Benchmark
  - # stress-ng --vm 1 --vm-bytes 1G --timeout 20s --metrics-brief

```
root@ubuntu:~# free -m
              total        used        free      shared  buffers  cached
Mem:       3952         123        3828          0         11        50
-/+ buffers/cache:       62        3890
Swap:          0          0          0
```

```
root@ubuntu:~# stress-ng --vm 1 --vm-bytes 1G --timeout 20s --metrics-brief
stress-ng: info: [1242] dispatching hogs: 1 vm
stress-ng: info: [1242] successful run completed in 20.04s
stress-ng: info: [1242] stressor      bogo ops real time  usr time  sys time  bogo ops/s  bogo ops/s
stress-ng: info: [1242]                                     (secs)  (secs)  (secs)  (real time) (usr+sys time)
stress-ng: info: [1242] vm                1250432    20.04    15.88     4.14   62403.02   62459.14
```

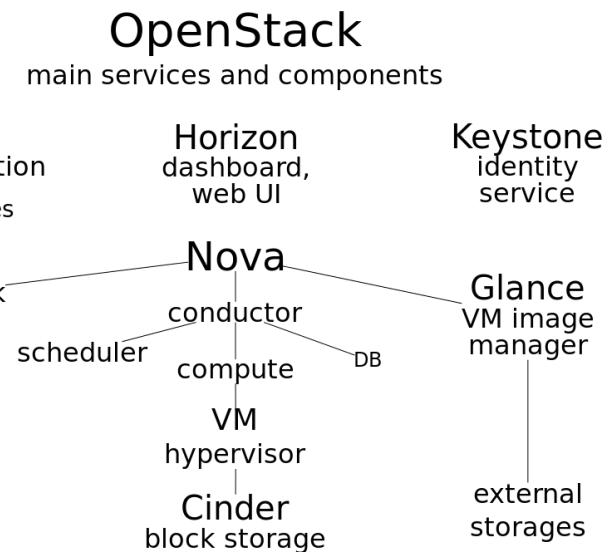
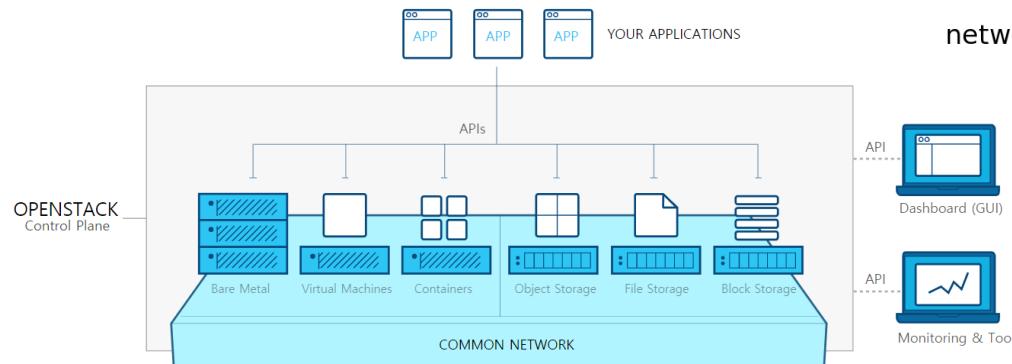
# LAB 12-2) OPENSTACK CLOUD PLATFORM

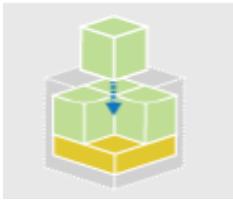
# OpenStack



openstack®

- OpenStack is a free and open-source software platform for cloud computing, mostly deployed as infrastructure-as-a-service (IaaS), whereby virtual servers and other resources are made available to customers.
- The software platform consists of interrelated components that control diverse, multi-vendor hardware pools of processing, storage, and networking resources throughout a data center.
- Users either manage it through a web-based dashboard, through command-line tools, or through RESTful web services.
- <http://www.openstack.org/>



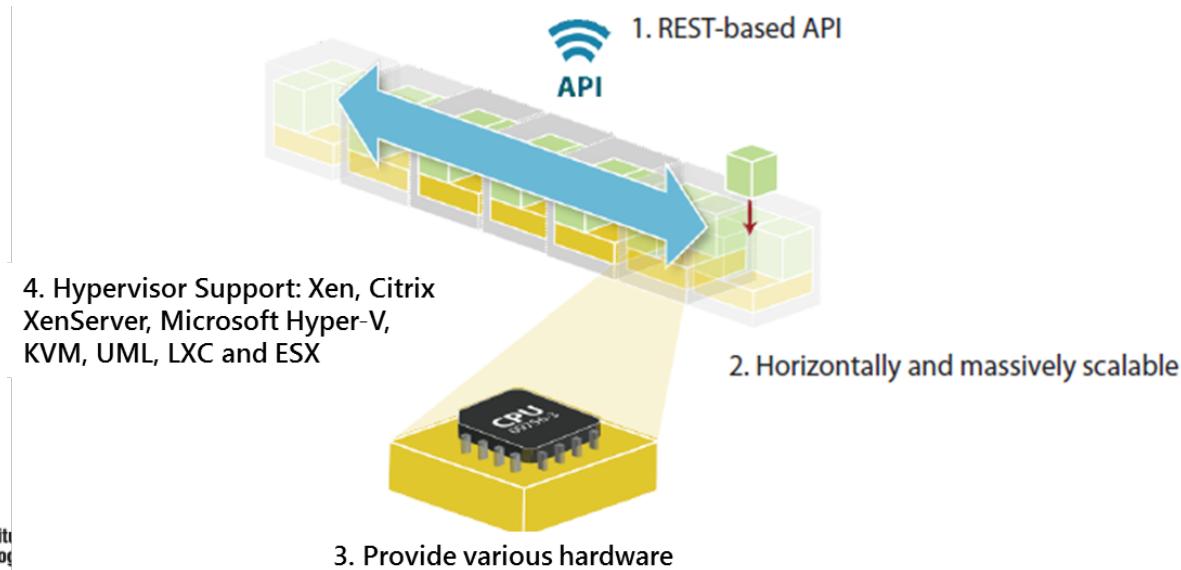


Referred from the article:

<http://ken.pepple.info/openstack/2012/09/25/openstack-folsom-architecture/>

# OpenStack Components (Nova)

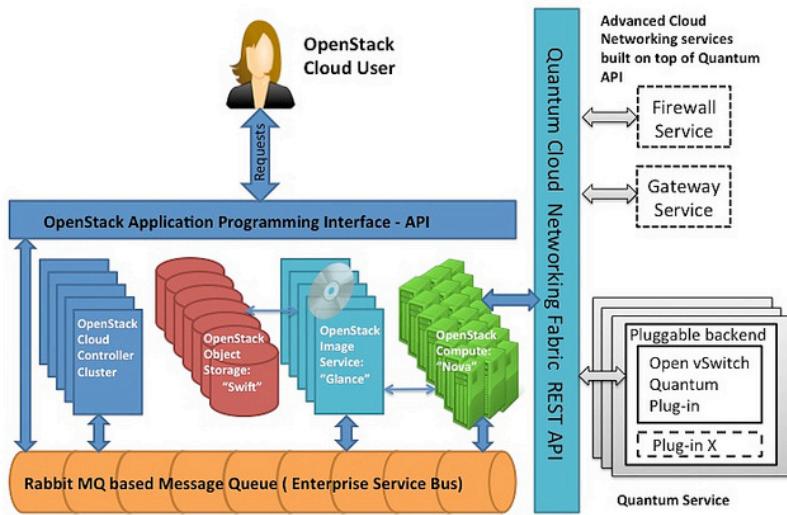
- Compute (code-named "**Nova**"):
  - Provides virtual servers upon demand
  - Manage and automate the large scale virtual compute instances pools using hypervisor
  - Work with widely available virtualization technologies, as well as bare metal and high-performance computing (HPC) configurations



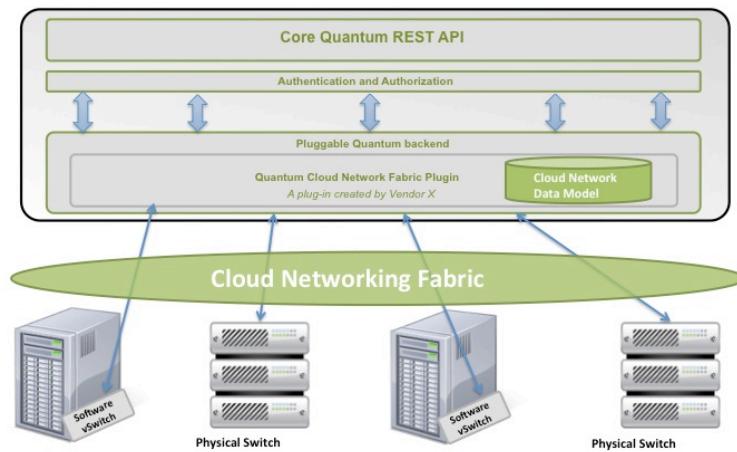


# OpenStack Components (Neutron)

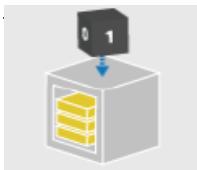
- Network (code-named "Neutron"):
  - “Network connectivity as a service” between interface devices(e.g., vNICs) managed by other OpenStack services
  - Application-level abstraction of networking that relies on plug-in implementations to map the abstraction(s) to reality
  - Quantum has a pluggable, scalable architecture to support many popular networking vendors and technologies
  - **Quantum is new in the Folsom release**



OpenStack + Quantum Integration Architecture

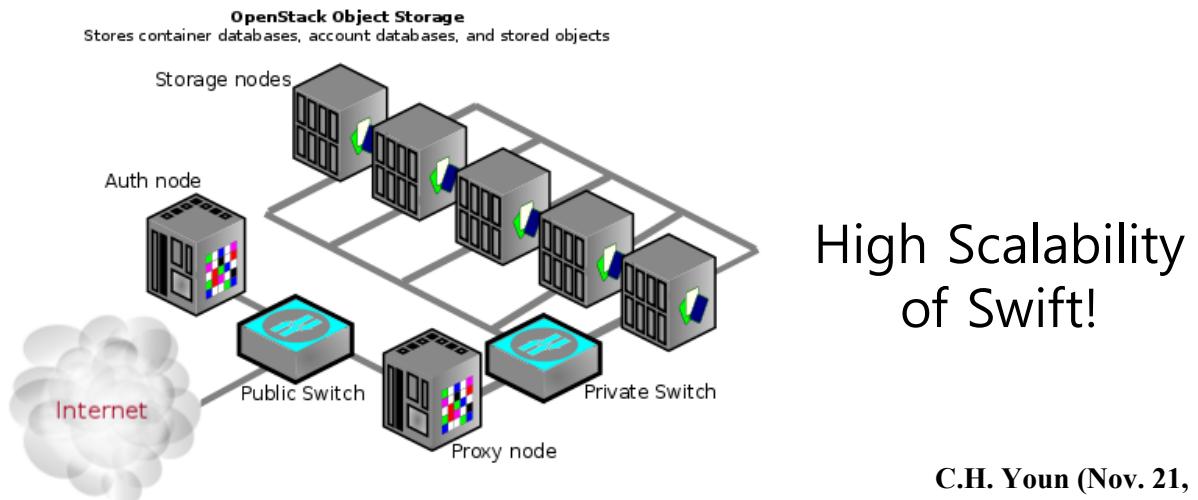


Reference Architecture – A Quantum Plugin



# OpenStack Components (Swift, Cinder)

- Object Store (code-named "**Swift**"):
  - Store or retrieve files like key-value data structure (Not general file system)
  - Open source software for implementing a reliable and large-scale cloud object storage service
  - Massively scalable redundant storage system
  - Objects and files are written to multiple disk drives spread throughout servers in the data center, with the OpenStack software responsible for ensuring data replication and integrity across the cluster
- Block Storage (code-named “**Cinder**”):
  - Provides persistent block level storage devices for use with OpenStack compute instances
  - Block storage system manages the creation, attaching and detaching of the block devices to servers





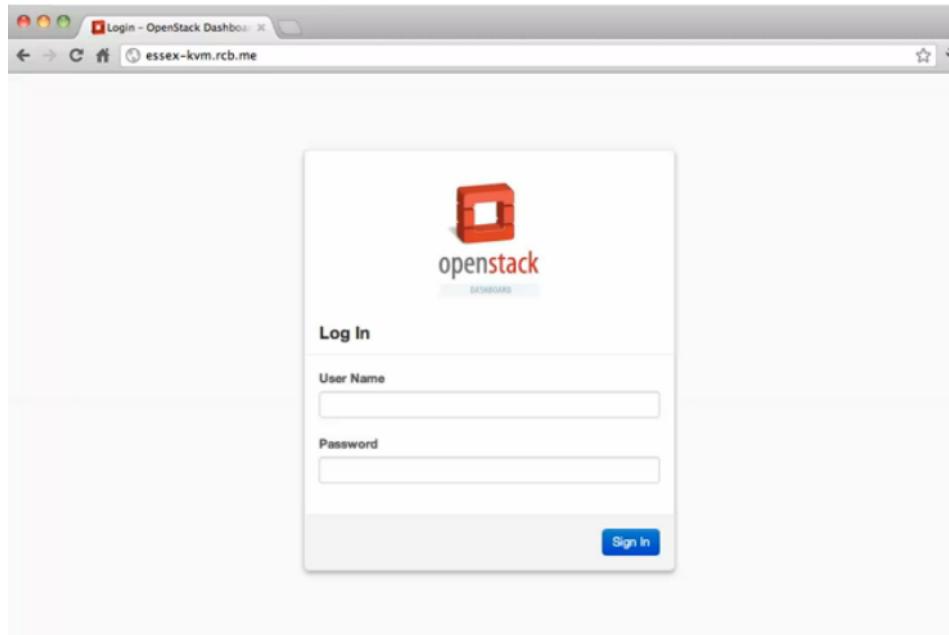
# OpenStack Components (Glance)

- Image (code-named "**Glance**"):
  - Provides a catalog and repository for virtual disk images
  - Discovery, registration and delivery services for disk and server images
  - Disk images are mostly commonly used for creating Virtual Machine instance in OpenStack Compute
  - Ability to copy or snapshot a server image and immediately store it away



# OpenStack Components (Keystone)

- Identity (code-named "**Keystone**"):
  - Provides authentication and authorization for all the OpenStack services
  - In API Service, keystone issues token to manage the session with clients
  - Provides a service catalog of services within a particular OpenStack Cloud



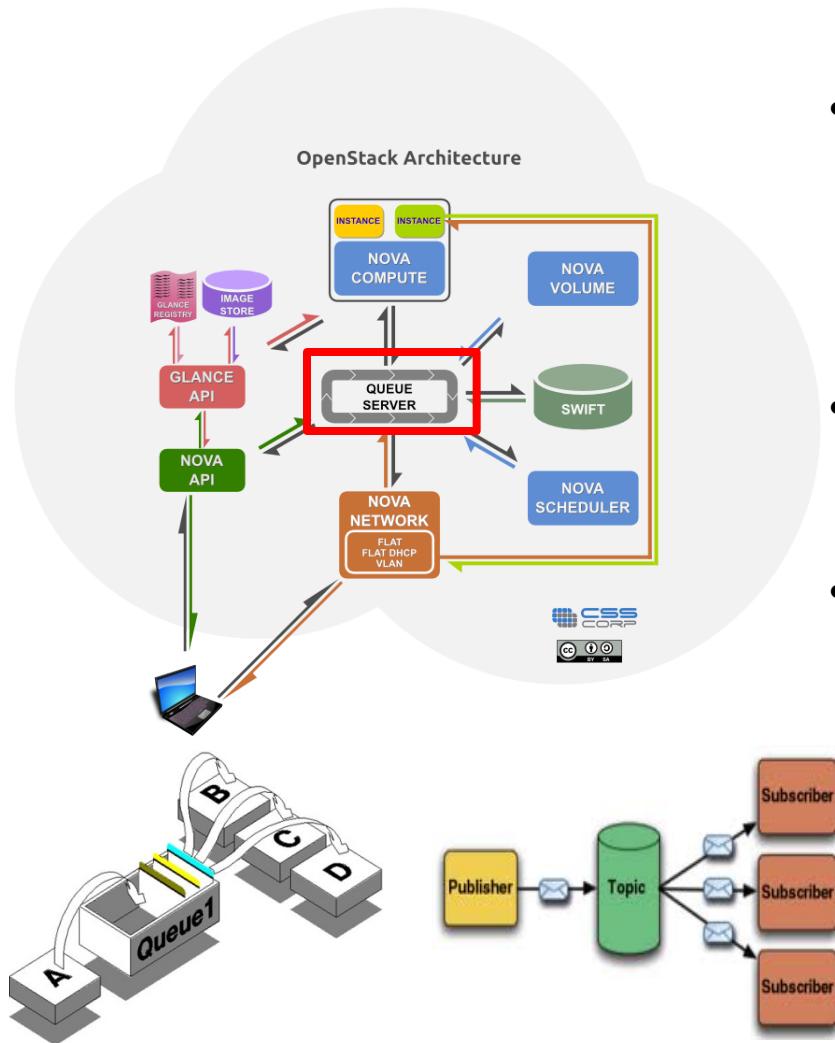


# OpenStack Components (Horizon)

- Dashboard (code-named "**Horizon**"):
  - A modular web-based user interface for all the OpenStack services.  
With web GUI
  - Perform most operations on cloud like launching an instance, assigning IP addresses and setting access controls

The screenshot shows the OpenStack Horizon dashboard at the URL [10.1.251.100/nova/instances\\_and\\_volumes/](http://10.1.251.100/nova/instances_and_volumes/). The dashboard has a sidebar on the left with links for Project, Manage Compute, Access & Security, Images & Snapshots, and Object Store. The main content area displays the "Instances & Volumes" page. A success message "Success: Instance \*test\* launched." is shown. The "Instances" section lists one item: "test" (512MB RAM | 1 VCPU | 0 Disk), Status: Build, Task: None, Power State: No State. Actions include Launch Instance, Terminate Instances, and Edit Instance. The "Volumes" section shows no items listed.

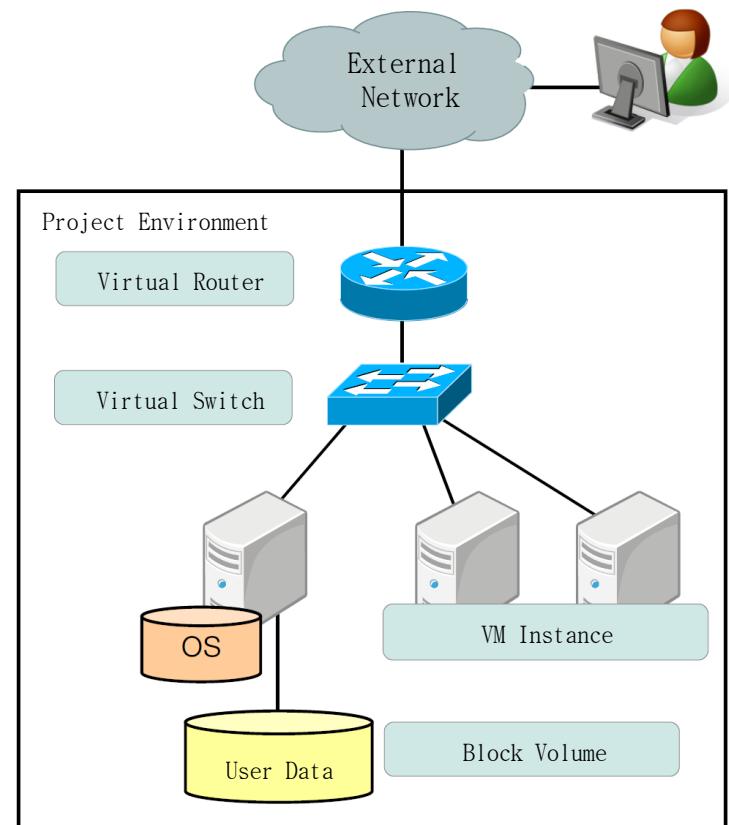
# RabbitMQ



- RabbitMQ is open source message broker software (i.e., message-oriented middleware, message queue system) that implements the Advanced Message Queuing Protocol(AMQP) standard
  - Software or hardware infrastructure supporting sending and receiving messages between distributed systems
  - Allows application modules to be distributed over heterogeneous platforms and reduces the complexity of developing applications that span multiple operating systems and network protocols
- In OpenStack, RabbitMQ mediates various modules and forwards messages from one to others
- Can download at <http://www.rabbitmq.com/>

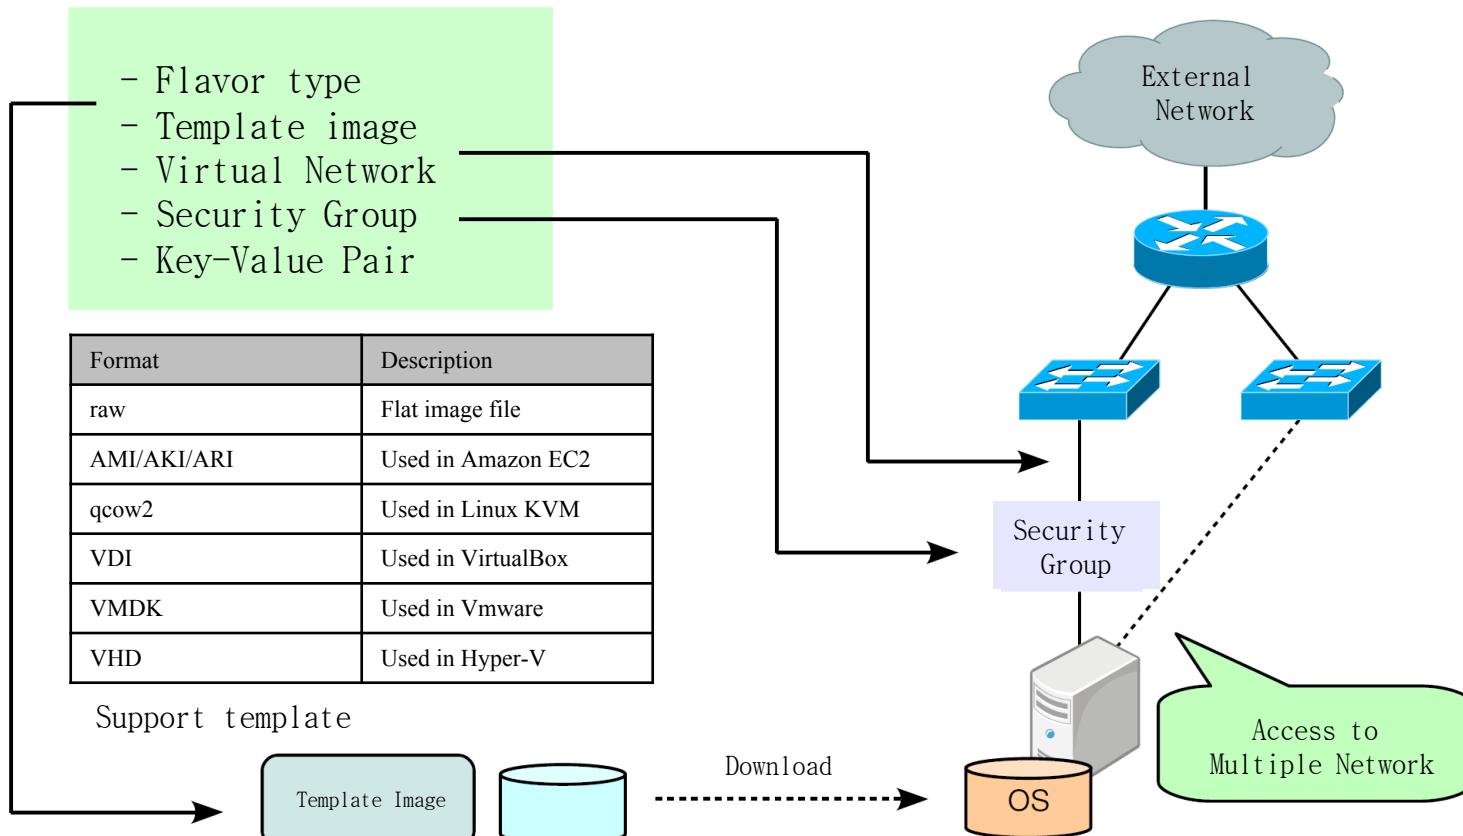
# Create VM Instance

- The end-user of the open stack can create and configure the following computing resources in the user's private tenants through a web console or API.
  - Virtual network
  - VM Instant
  - Block volume
- Each user belongs to some projects.
  - Users in the same project share common computing resources in the project environment.
  - Each project owns (virtual) computing resources that are independent of other projects



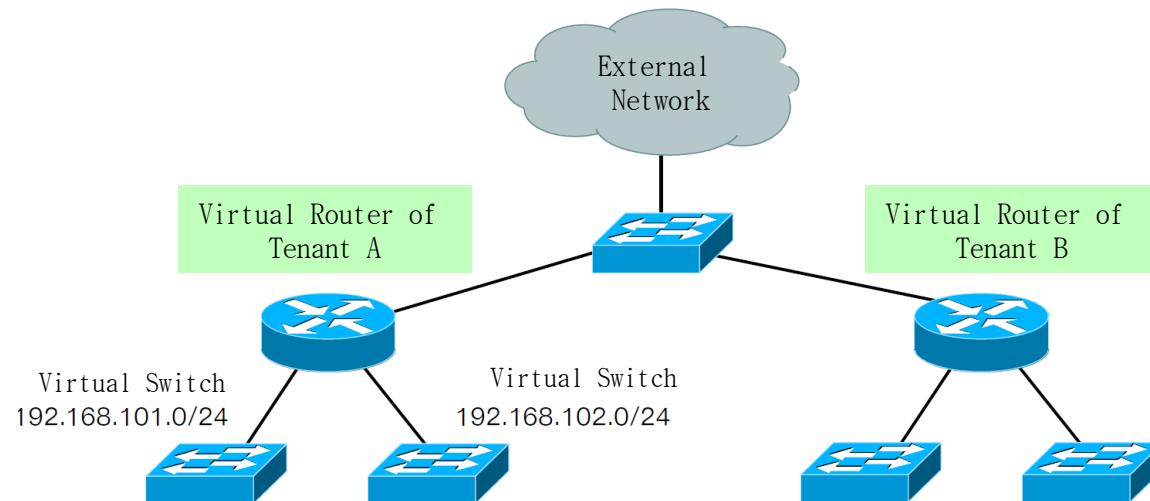
# Create VM Instance

- When starting a new VM instance, you must specify the following options:



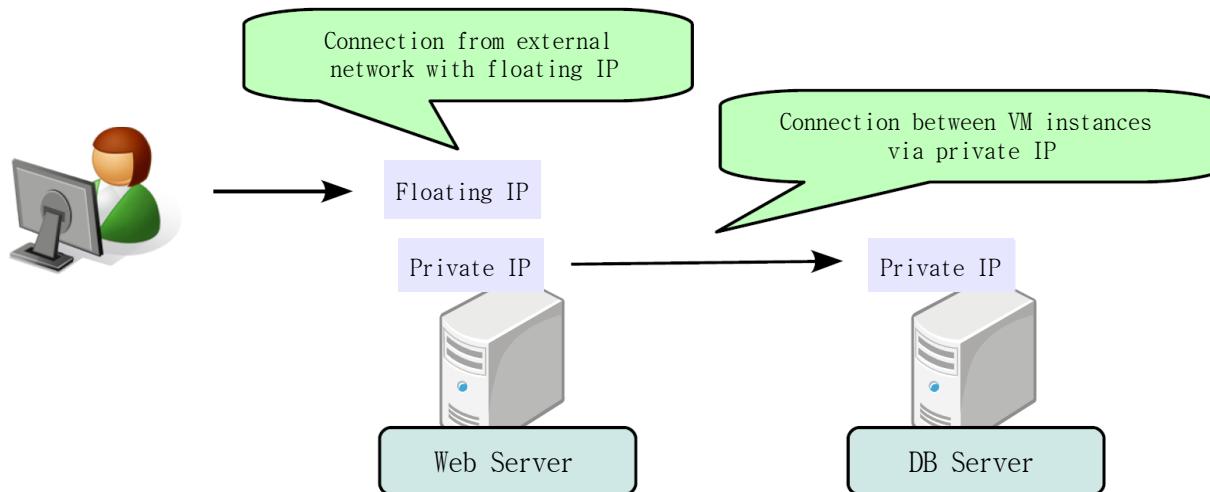
# Logical View of Open Stack Virtual Networks

- Tenant users add virtual switches to the bottom of the router and assign a private subnet address. Subnets can be used in conjunction with other tenants.
- When starting the instance, the end user selects the virtual switch to connect to.
- The number of virtual NICs in the instance corresponds to the number of switches to connect. It assigns private IPs through DHCP.



# Private IP and Floating IP

- When accessing from an external network, assign "Floating IP" to the VM instance.
  - IP addresses of external networks that can be used as floating IP are pooled in advance and distributed to each tenant.
  - Floating IP is NAT-converted from the virtual router to the corresponding private IP.
  - Accessing the external network from the VM instance is possible without floating IP allocation. In this case, the IP Masquerade function of the virtual router is used.

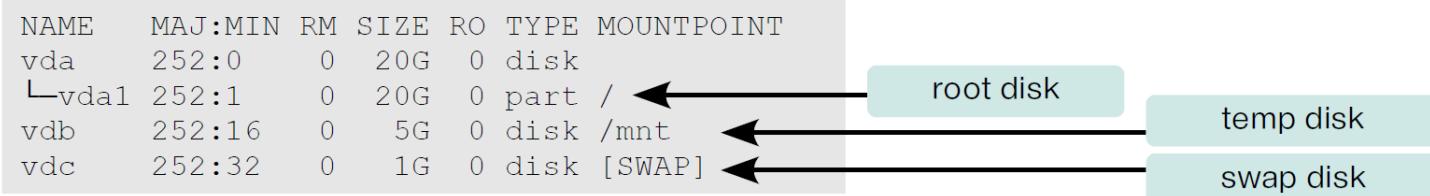


# The instance type and disk space

- The following is the default instance type.
  - The root disk is copied from the template image and expanded to the specified size (except m1.tiny).

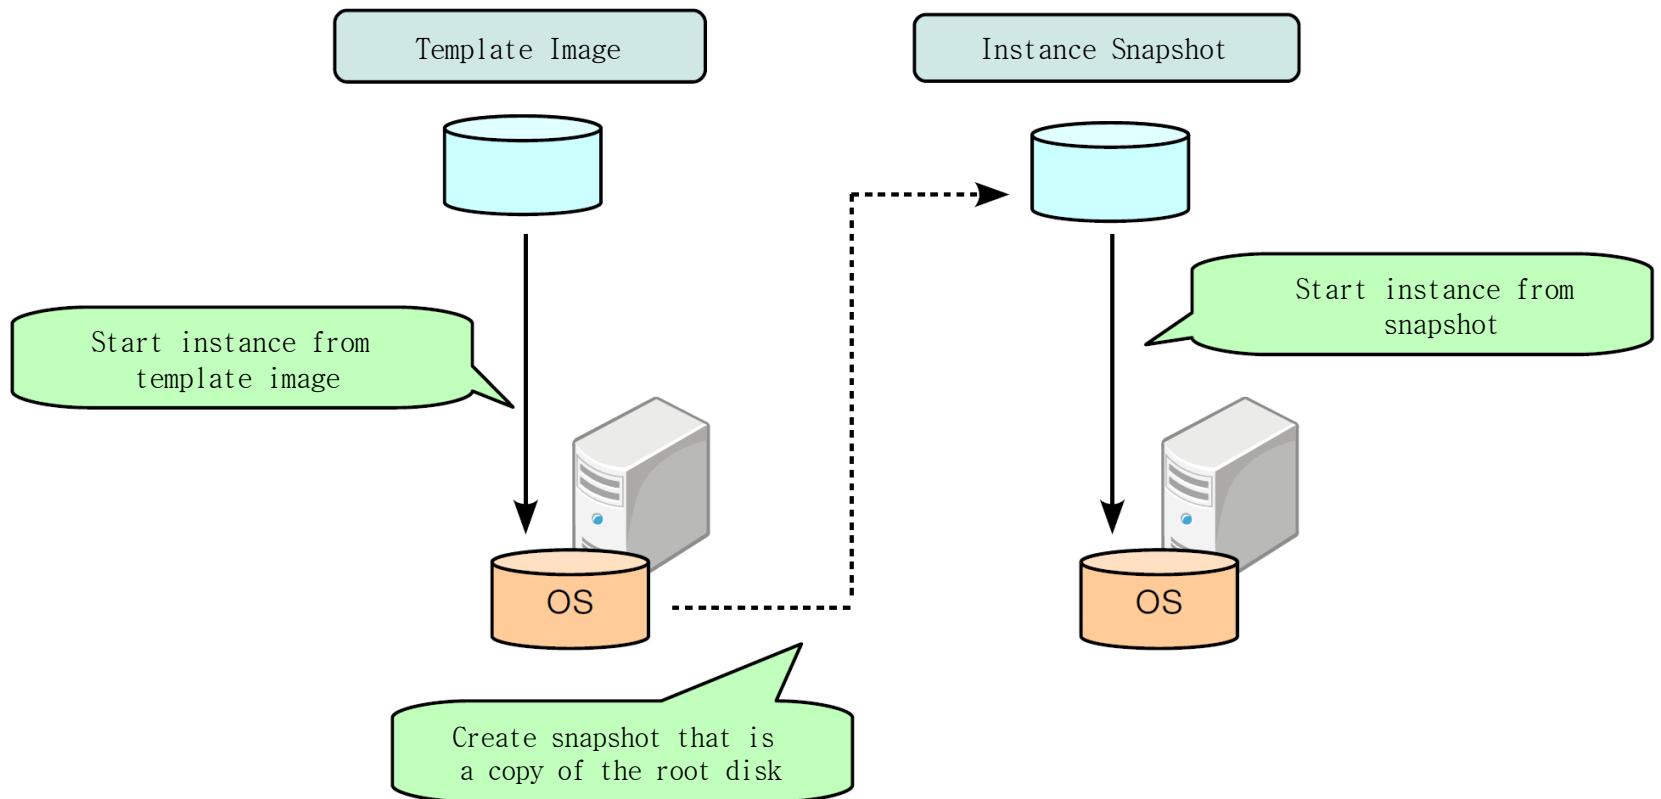
Flavor type	vCPU	Memory	Root disk	Temp disk	Swat disk
m1.tiny	1	512MB	0GB	0	0
m1.small	1	2GB	20GB	0	0
m1.medium	2	4GB	40GB	0	0
m1.large	4	4GB	80GB	0	0
m1.xlarge	8	8GB	160GB	0	0

- Administrators can define new instance types.
  - The following is an example using temp disk and swap disk.
  - Persistent data is stored elsewhere (typically block volumes) because these disks are deleted when the instance is destroyed.



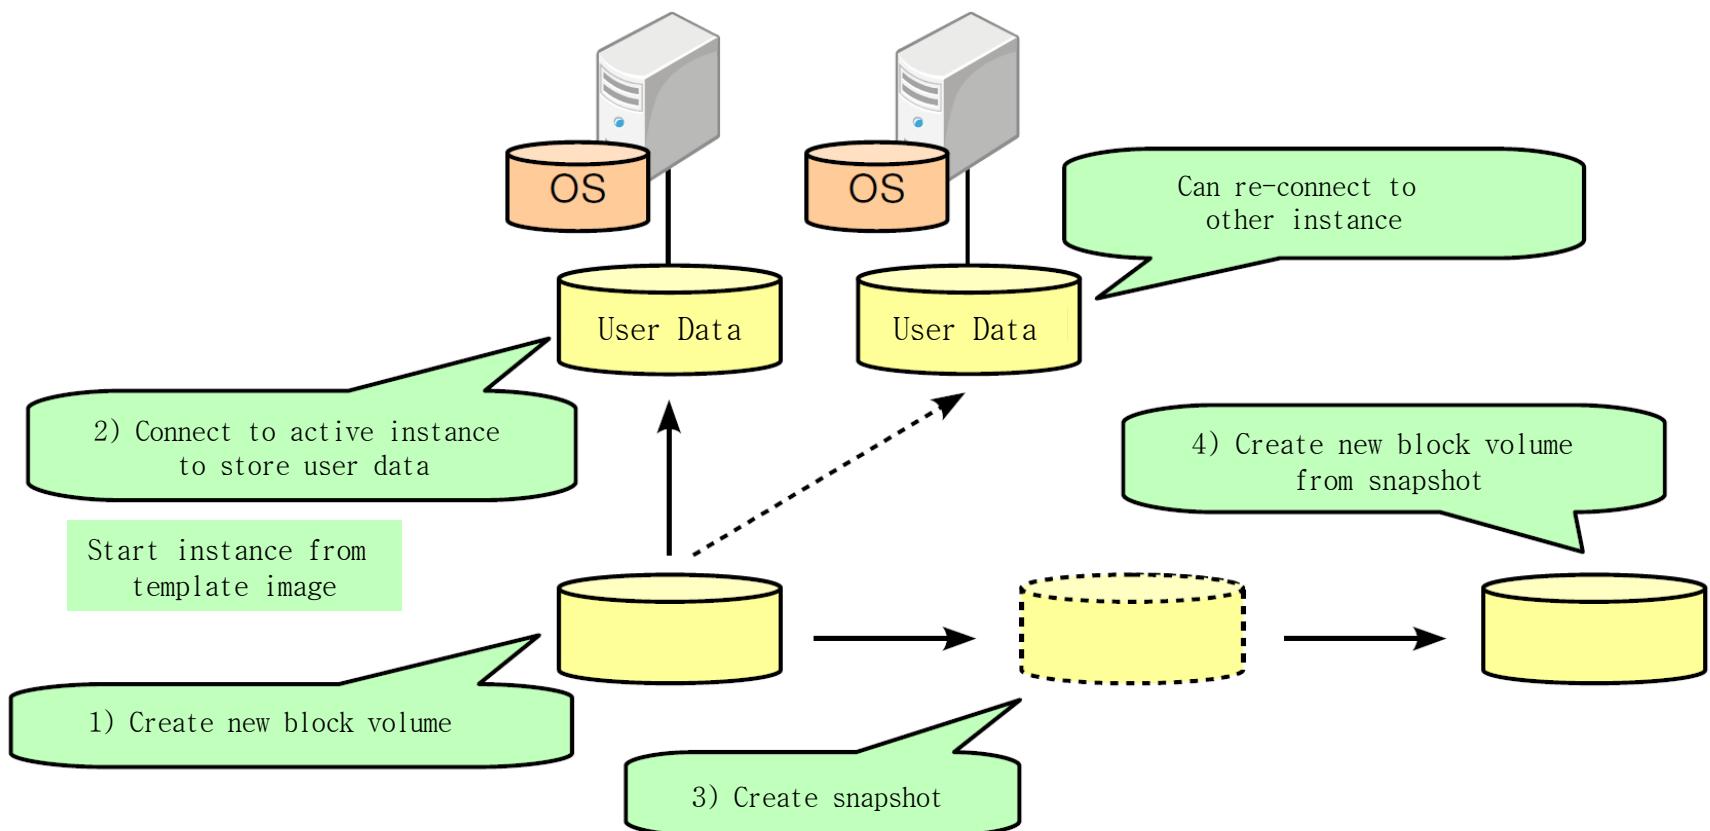
# Snapshot of VM instance

- You can create a snapshot of the running instance, copy the root disk, and reuse it as a template image.



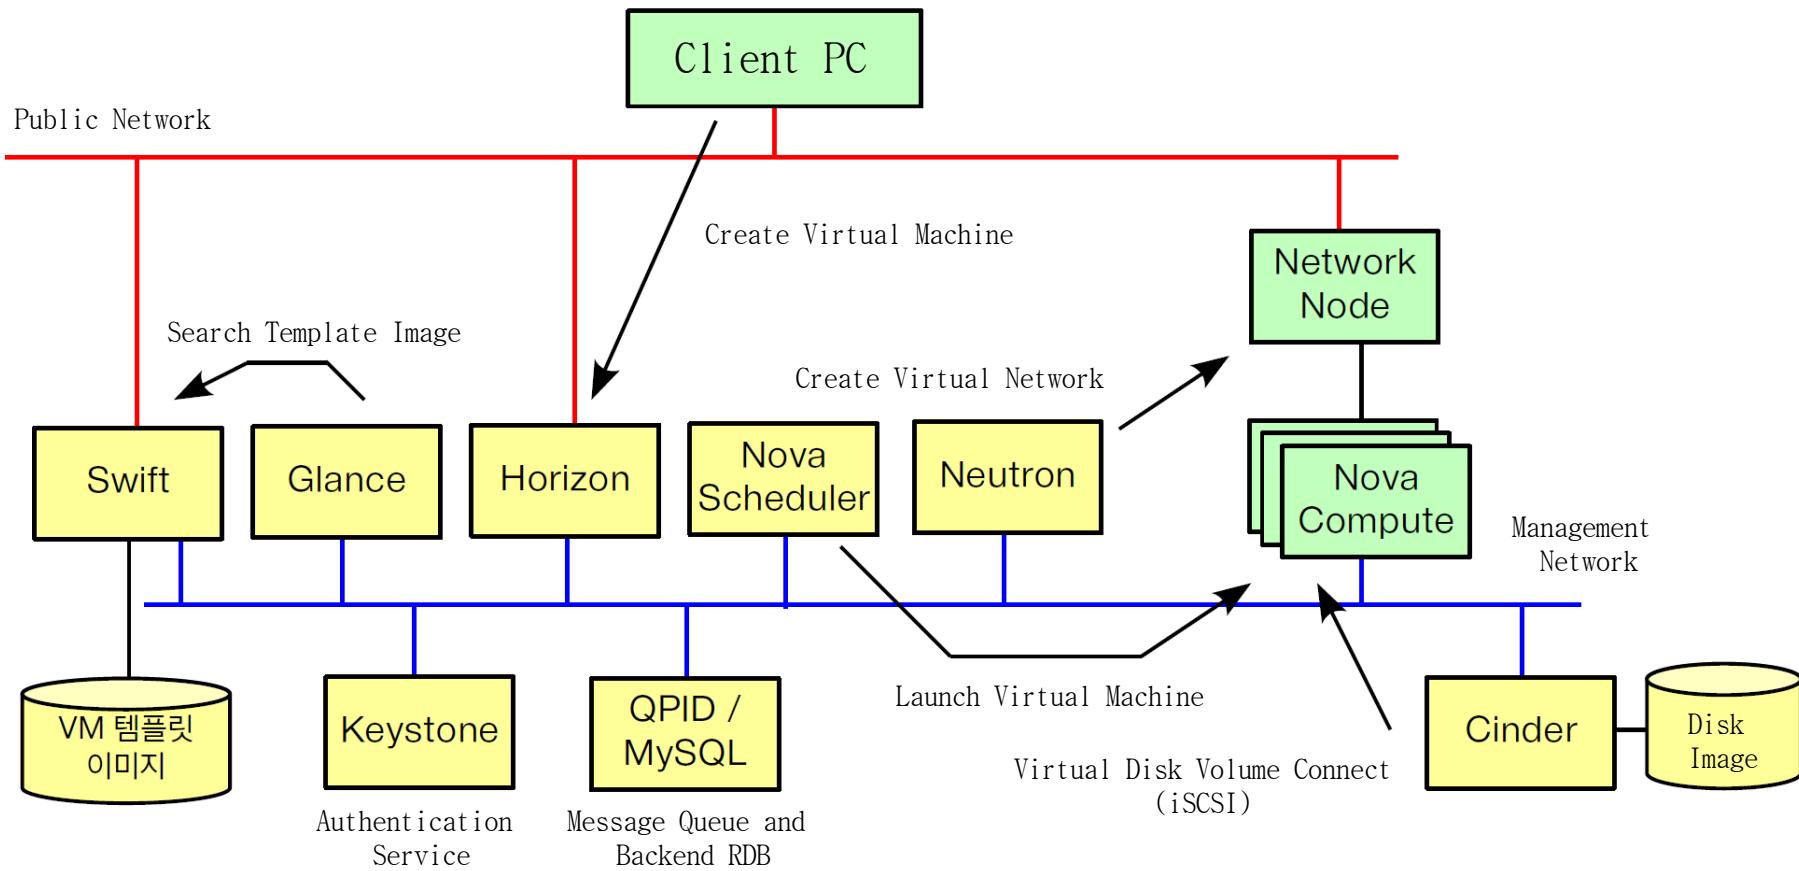
# Persistent data store, block volume

- The block volume is not deleted when the VM instance is removed. It can be used as a persistent data store.



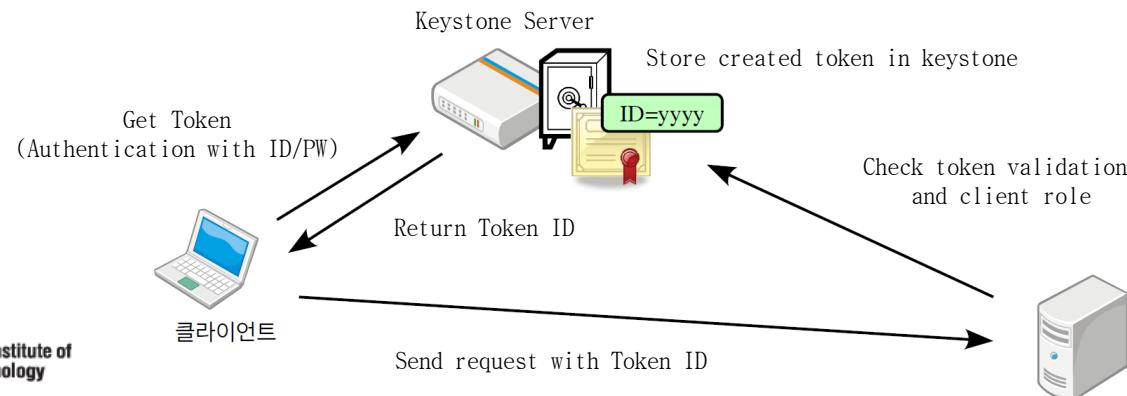
# System Operation

- Modules work together through REST API calls and message queues.
  - Operation can be automated by external program through REST API.

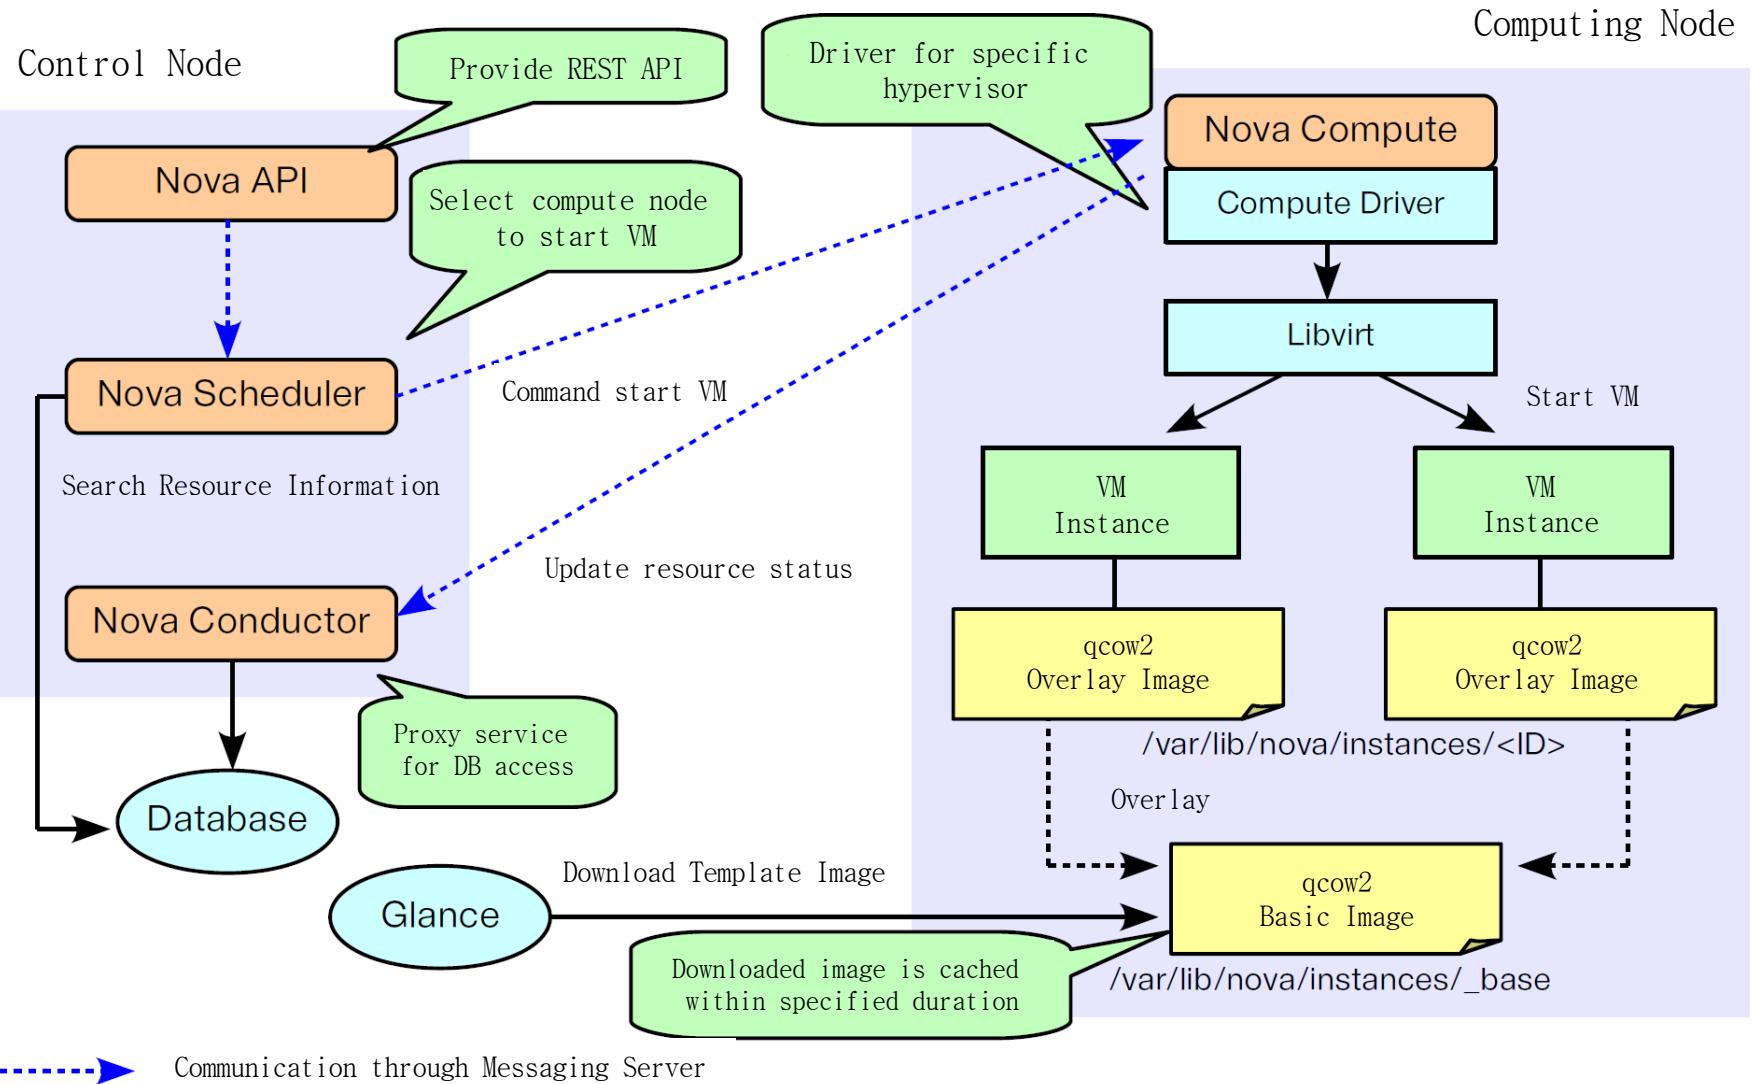


# User authentication for API requests

- Authentication is required before sending a request to the API.
  - The end user / component obtains a "token" for the API operation from the keystone before sending the request to the API. (The user ID of each component is in the Keystone.)
  - Once the token is obtained, the URL for the target API is retrieved from the keystone. The end user must know only the URL to the Keystone API in advance.
- Since the open stack client makes many API calls to various components,
- Authentication using ID / password is not desirable in terms of security and performance.
- Instead, the client obtains a "token" as a "license" for the API call in advance,
- Sends the token ID to the component to be used.
  - The component receiving the request validates the token ID in the Keystone before approving the request.
  - The generated token is stored in the Keystone for a period of time (default: 24 hours). When the client expires
  - The token can be used again. The client does not need to obtain a token for each request invocation.

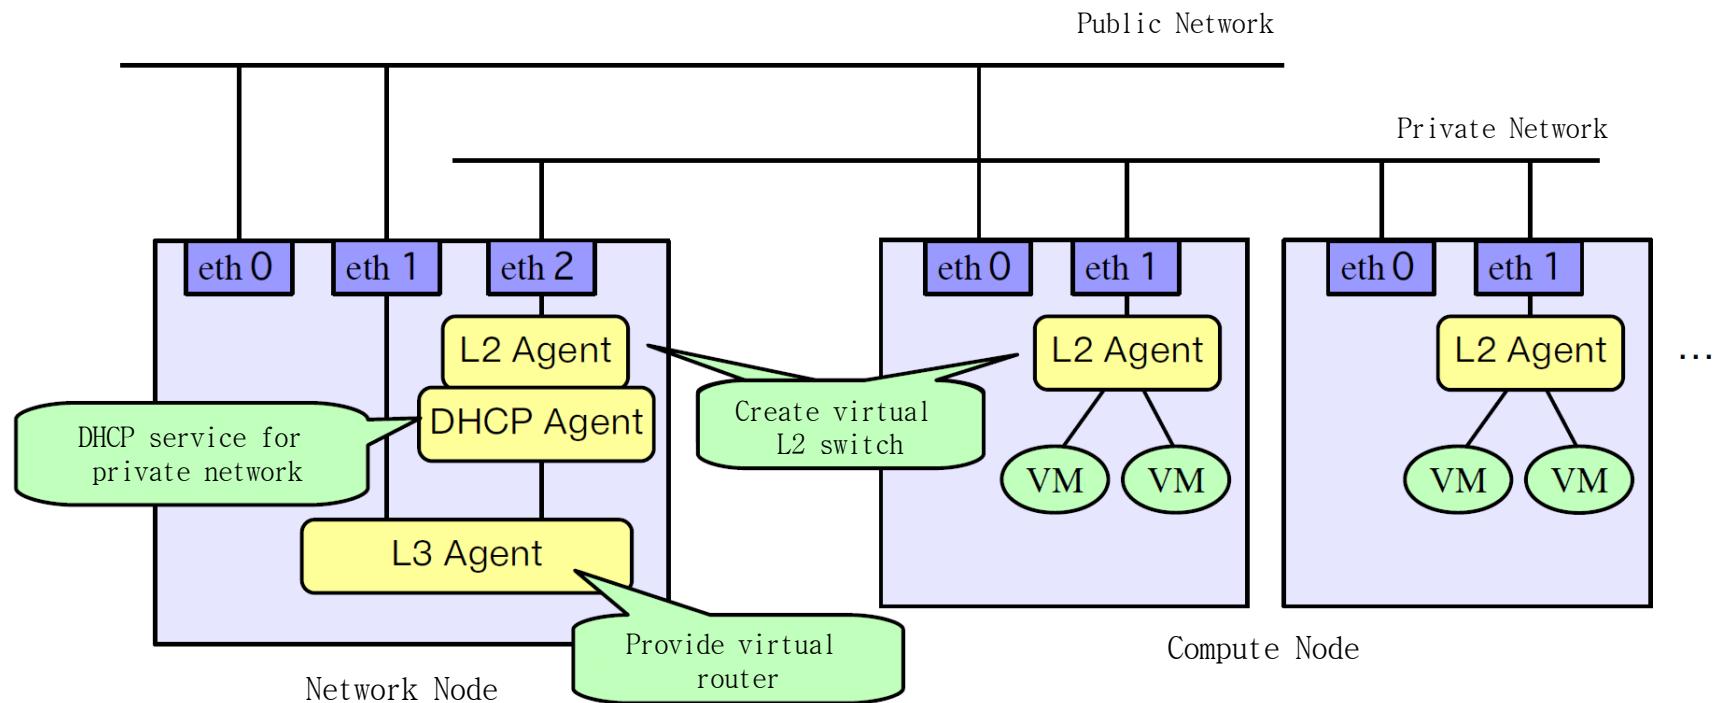


# Nova Internal



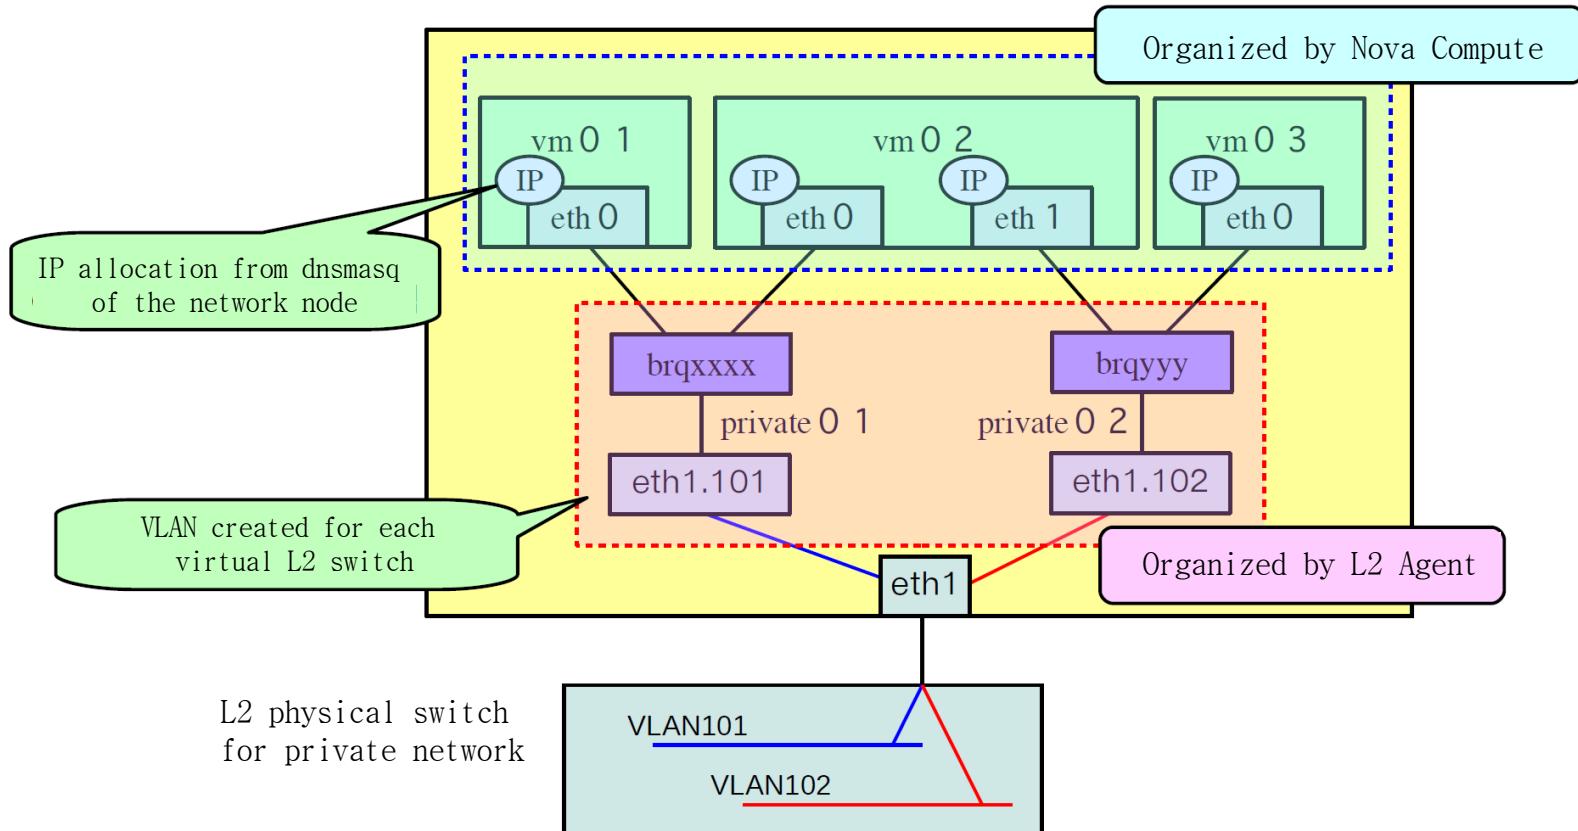
# Network Composition

- Here is an example of a typical configuration using the LinuxBridge plugin or the Open vSwitch plugin:
  - The L3 Agent of a network node provides a virtual router function to connect private and public networks. (“eth0” on each node is used to access host Linux, not on VM instance communication)

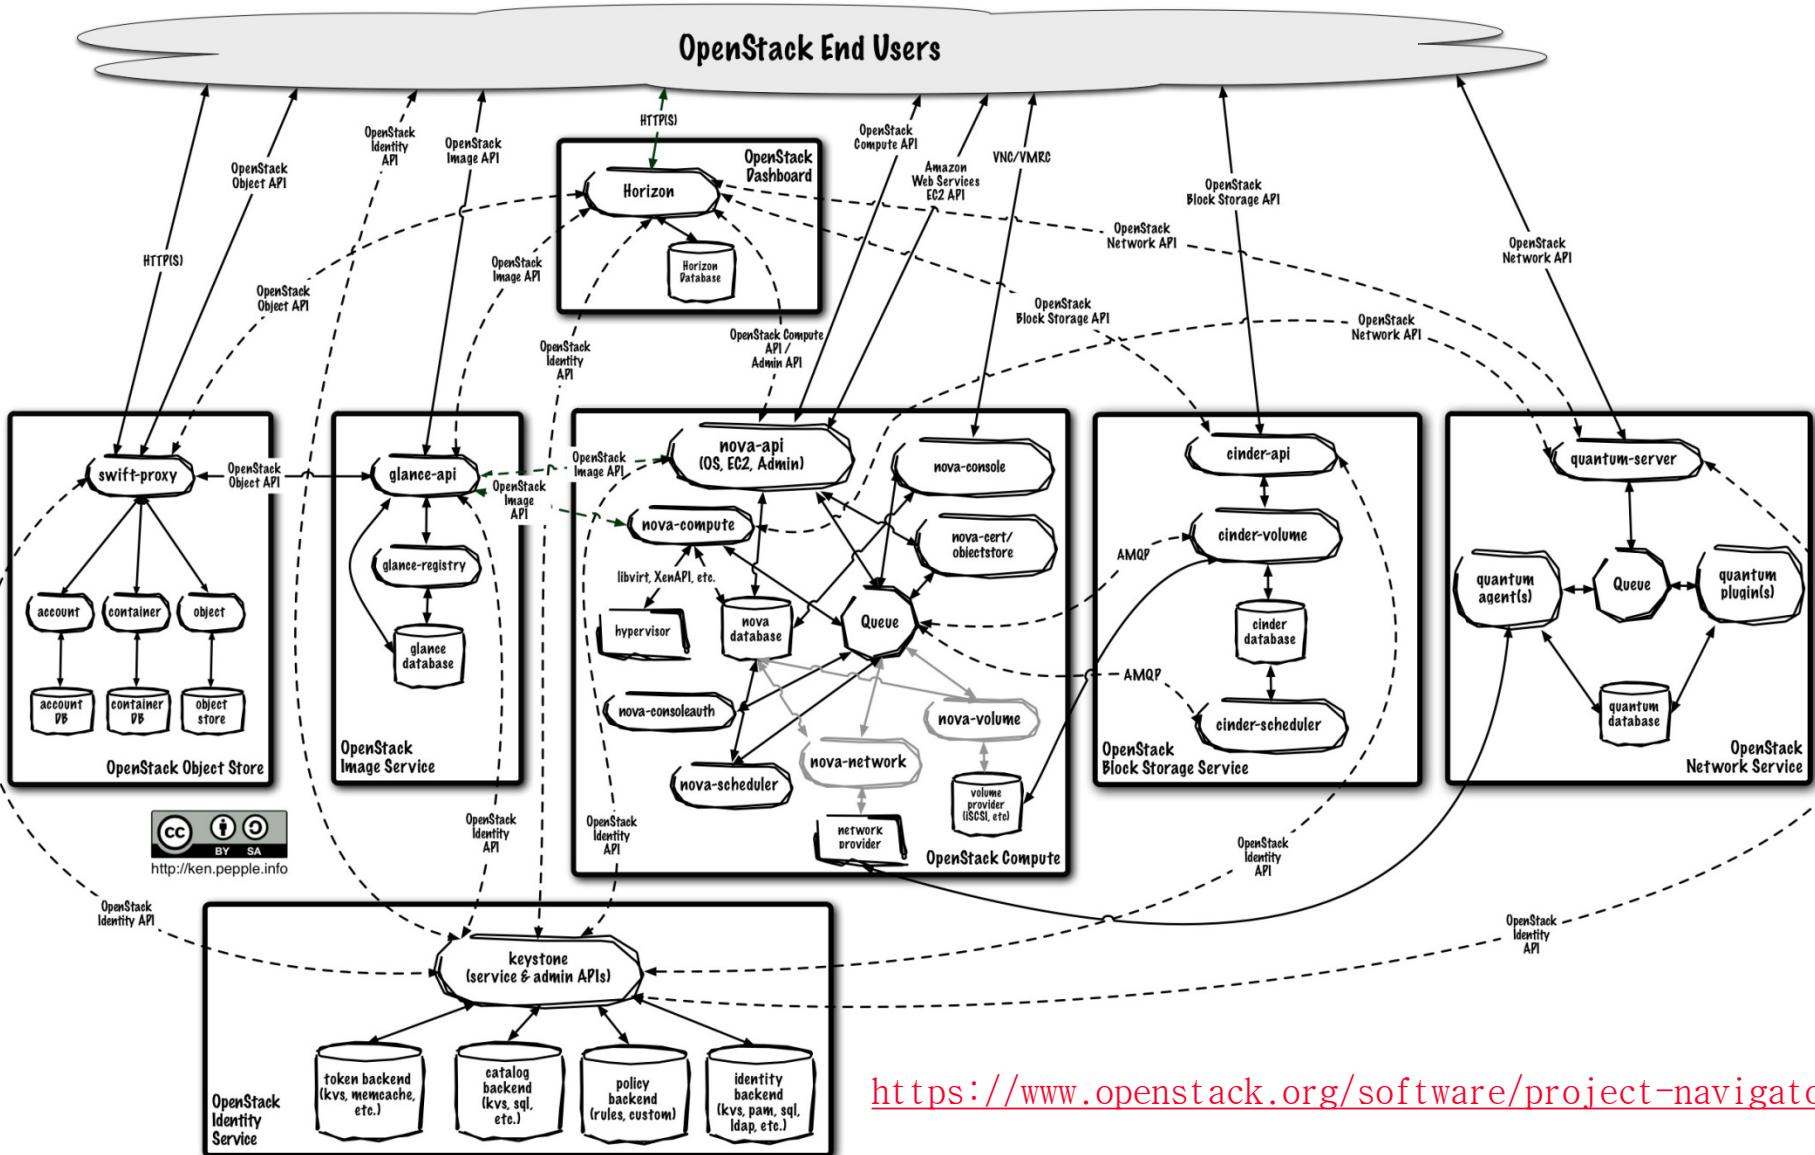


# Compute Node Internal

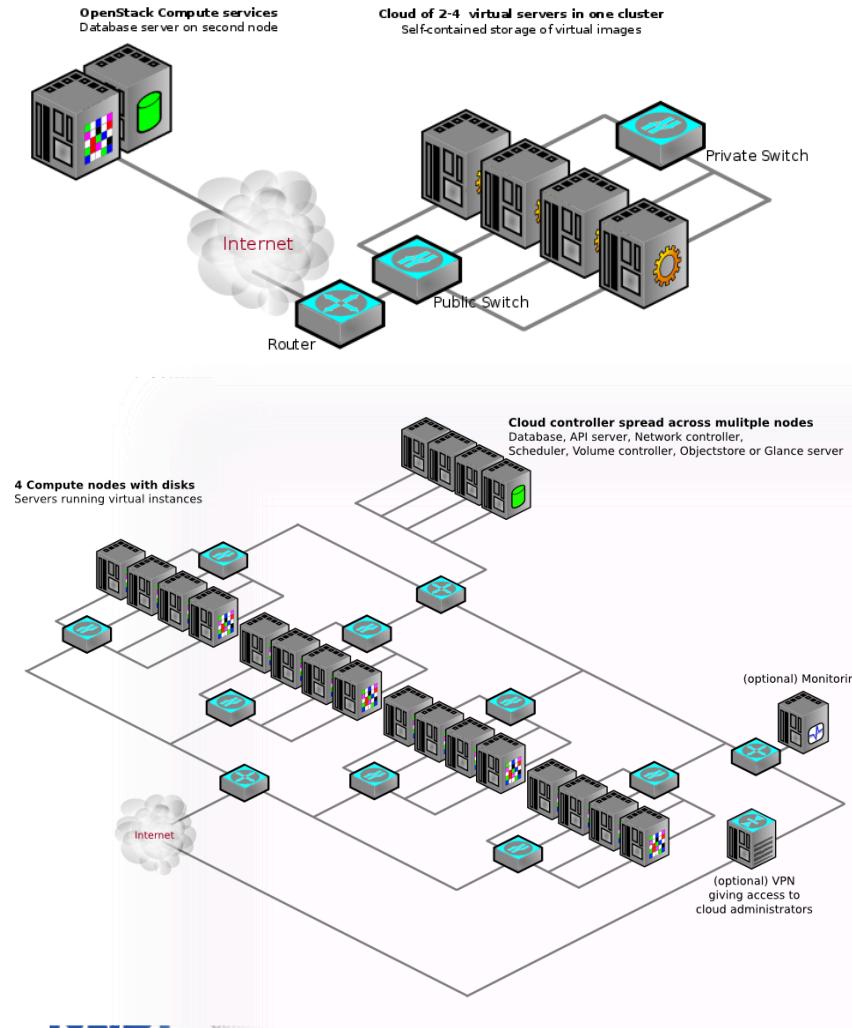
- A Linux bridge is created for each virtual switch. The network traffic of each switch other than the Compute node is separated into VLANs.



# Detail OpenStack Functions



# Examples of Installation Architectures

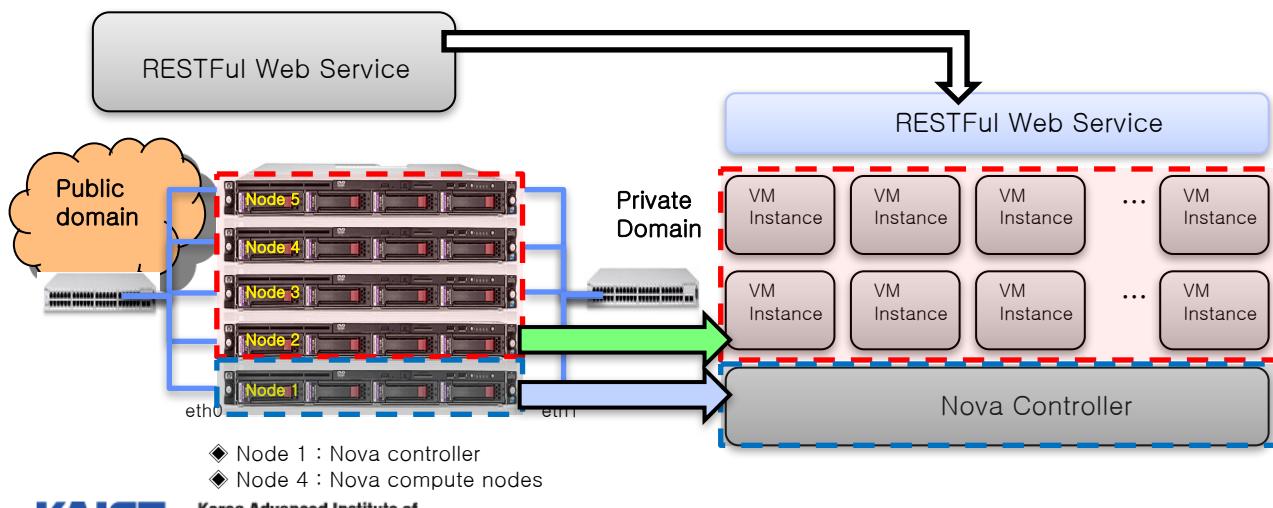


- Single node: Only one server runs all nova-services and also drives all the virtual instances. Use this configuration only for trying out OpenStack Compute, or for development purposes.
- Two nodes: A cloud controller node runs the nova- services except for **nova-compute**, and a compute node runs **nova-compute**. A client computer is likely needed to bundle images and interfacing to the servers, but a client is not required. Use this configuration for proof of concepts or development environments.
- Multiple nodes: You can add more compute nodes to the two node installation by simply installing **nova-compute** on an additional server and copying a nova.conf file to the added node. This would result in a multiple node installation. You can also add a volume controller and a network controller as additional nodes in a more complex multiple node installation. A minimum of 4 nodes is best for running multiple virtual instances that require a lot of processing power.

# Installation Environment

- Rack Mounted Cluster Machines  
at E3-2 #3211

OpenStack Platform	
<b>Hypervisor</b>	KVM
<b>H/W Specification</b>	Intel Xeon E5620 2.40GHz, Core 16, MEM 16G, HDD 1T, 5 Node
<b>S/W Specification</b>	OS: Ubuntu 14.04
<b>IP addresses</b>	Eth0 (143.248.152.61~143.248.152.65)



# OpenStack Installation

- To install OpenStack in basic way, we have to follow these documentations:  
<https://docs.openstack.org/install-guide/>
  - You should follow about 200 pages instructions and it will be complex and take long-time
  - So we will use linux shall script based installation supporter DevStack (<http://devstack.org/>)
- DevStack needs minimum configuration for building OpenStack platform (but this property become weakness in commercialization build)

# OpenStack Installation

## 1. Prepare the Ubuntu Environment in 16.04.3 LTS Desktop 64 bit Version

- You can use VMWare or VirtualBox but not recommended (OpenStack requires minimum 1.5GB memory)

## 2. Add Stack User

- You can create a separate stack user to run DevStack with

```
$ sudo useradd -s /bin/bash -d /opt/stack -m stack
```

- Since this user will be making many changes to your system, it should have sudo privileges:

```
$ echo "stack ALL=(ALL) NOPASSWD: ALL" | sudo tee /etc/sudoers.d/stack  
$ sudo su - stack
```

## 3. Download DevStack

```
$ git clone https://git.openstack.org/openstack-dev/devstack  
$ cd devstack
```

# OpenStack Installation

## 4. Create a local.conf file

- Specification for installation
- Create local.rc file with following contents

```
[[local|localrc]]  
HOST_IP=192.168.122.1  
ADMIN_PASSWORD=secret  
DATABASE_PASSWORD=$ADMIN_PASSWORD  
RABBIT_PASSWORD=$ADMIN_PASSWORD  
SERVICE_PASSWORD=$ADMIN_PASSWORD
```

## 5. Start the devstack installation

- Installation takes very long time (about 40 minutes)

```
./stack.sh
```

# Specific configuration file and shall script files

localrc

```
sudo iptables -A POSTROUTING -t mangle -p udp --dport
  68 -j CHECKSUM --checksum-fill
tools/upload_image.sh
http://ancl.kaist.ac.kr/data/SOC2013/openstack/SOC.im
g
```

```
# network
FLAT_INTERFACE=br100
FIXED_RANGE=192.168.2.1/24
FIXED_NETWORK_SIZE=256

# vnc
VNCSERVER_LISTEN=0.0.0.0
VNCSERVER_PROXYCLIENT_ADDRESS=$HOST_IP

# logs
DEST=/opt/stack
LOGFILE=$DEST/logs/stack.sh.log
SCREEN_LOGDIR=$DEST/logs/screen

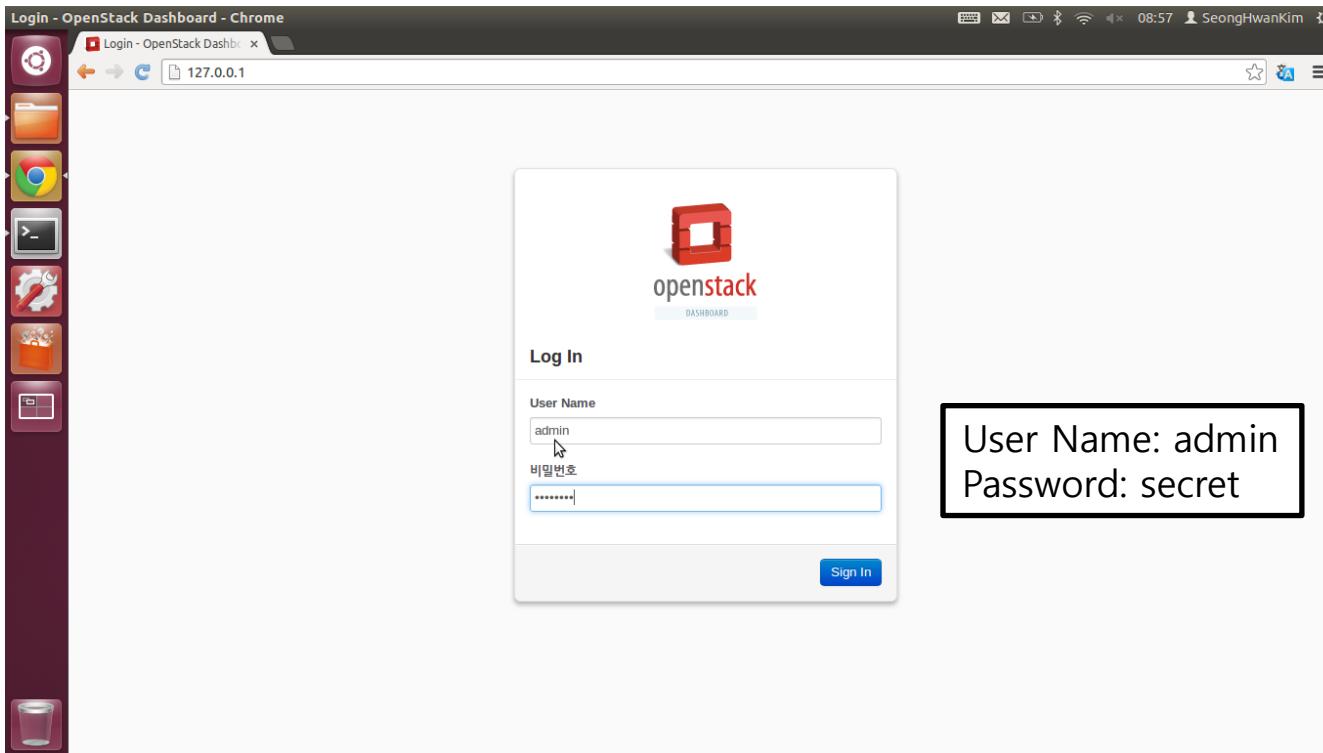
# system password
ADMIN_PASSWORD=ee614soc
MYSQL_PASSWORD=ee614soc
RABBIT_PASSWORD=ee614soc
SERVICE_PASSWORD=ee614soc
SERVICE_TOKEN=socservicetoken

# cinder
VOLUME_GROUP="cinder-volume"
VOLUME_NAME_PREFIX="volume-"
```

- \* When installation meet unexpected error or exit,  
please run /home/stack/devstack/unstack.sh and clean.sh

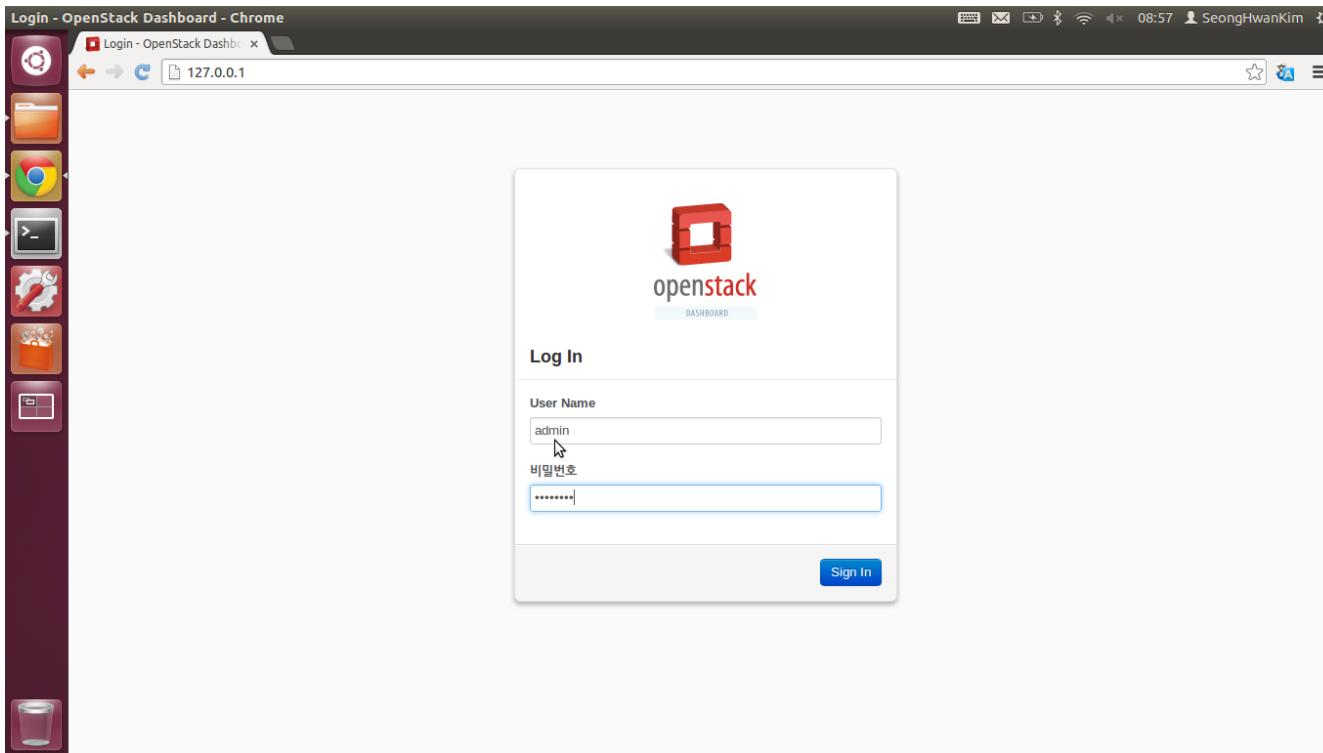
# OpenStack Installation Verification

1. Open the Web Browser and access to <http://127.0.0.1>
  - Authenticate access with user account: ‘admin’, password: ‘secret’



# OpenStack Installation Verification

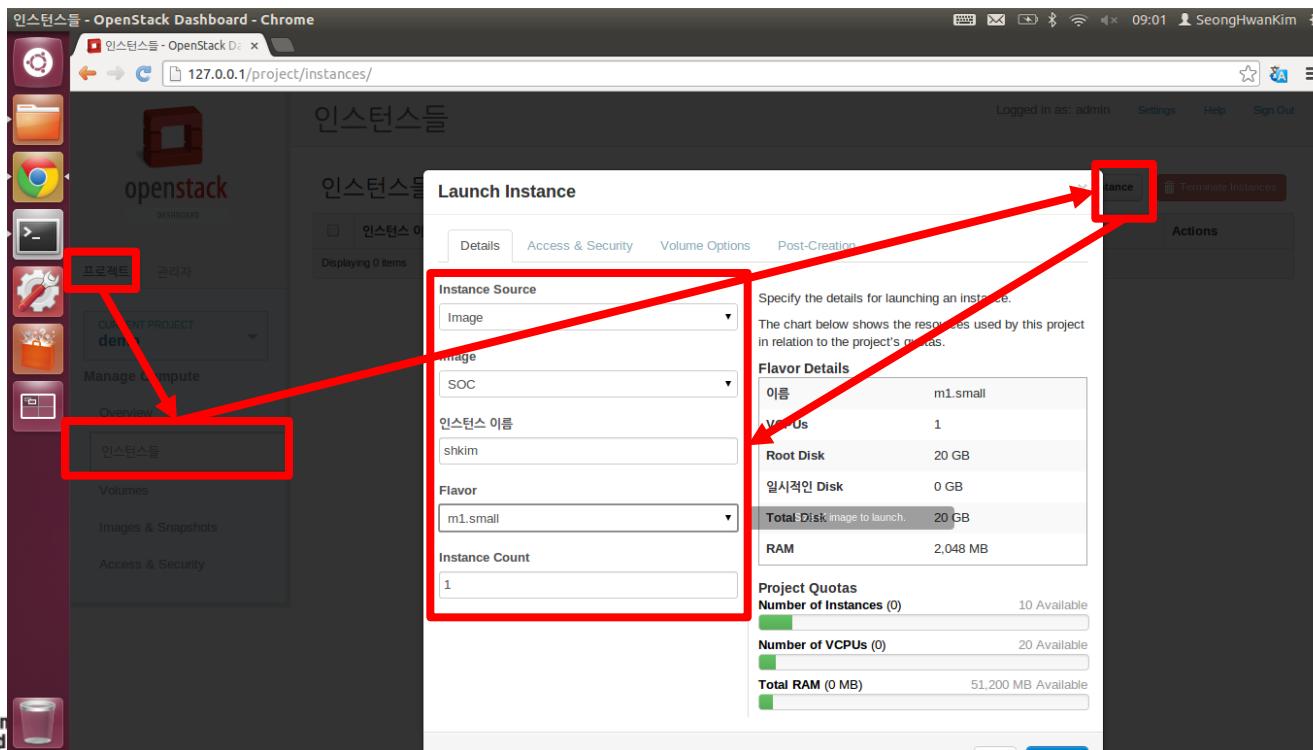
1. Open the Web Browser and access to  
<http://143.248.146.60/horizon>



# OpenStack Installation

## Verification

2. Access ‘project→instances→Launch Instance’ to create instance
  - Create instance with image=‘SOC’, instance name=‘any name’, Flavor=‘m1.small’, instance count=‘1’

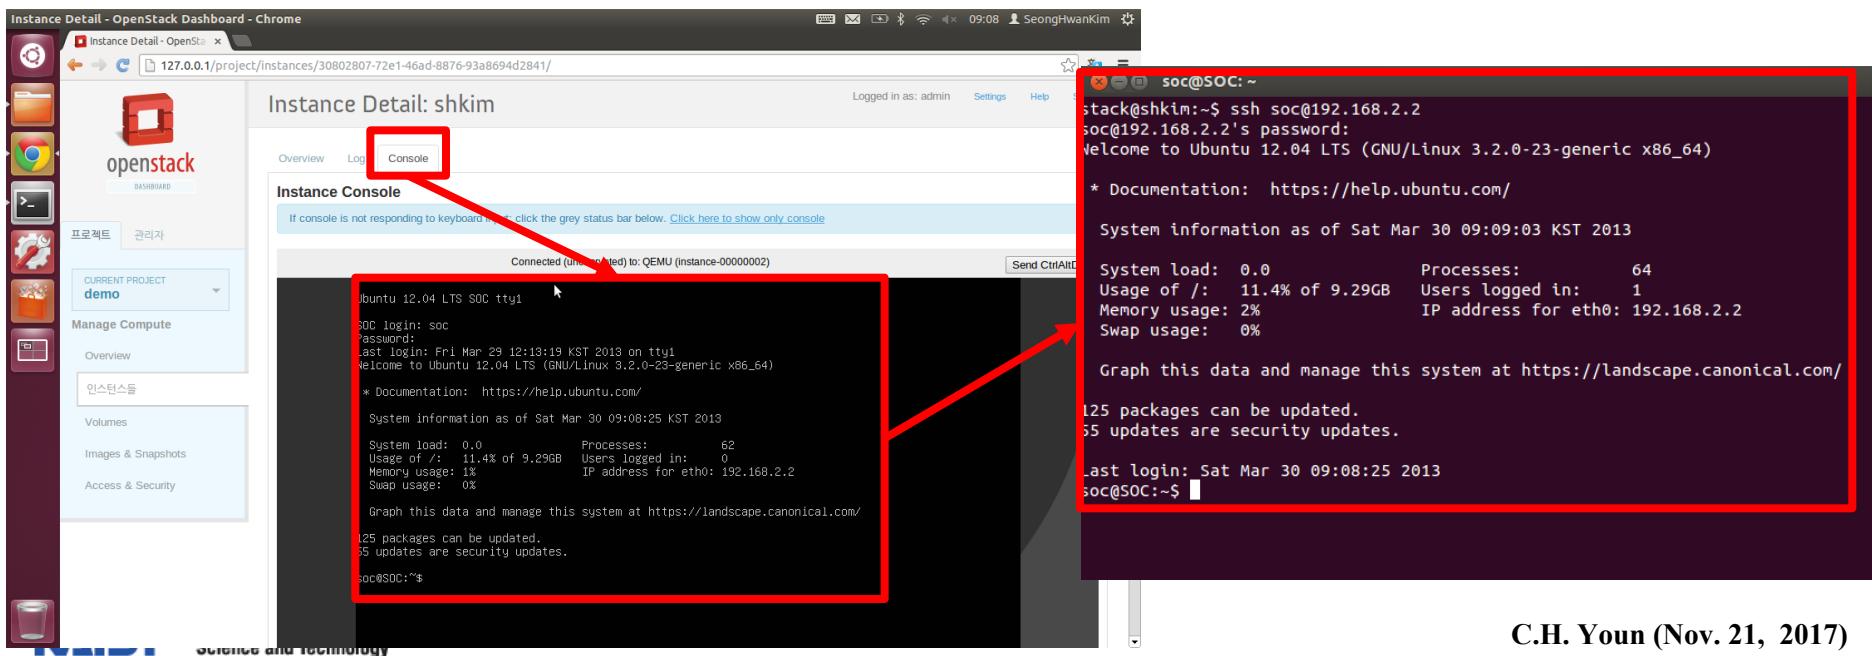


# OpenStack Installation

## Verification

3. When instance status becomes to ‘Active’, try to access VM

- Click the instance name and access to the VM with console label (VNC Console)
- Open the terminal and access VM with ssh protocol  
→ \$ ssh soc@192.168.2.2
- VM User Account: soc, Password: ee614soc



# OpenStack Control

- Please watch the movie clip in  
<http://www.openstack.org/software/openstack-dashboard/>  
and be familiar with OpenStack Platform and Dashboard

